

Calibrated Model Criticism Using Split Predictive Checks

Jonathan Huggins
Boston University

with Jiawei Li

Model checking: can we have it all?

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient
 3. well-calibrated

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient
 3. well-calibrated
 4. well-powered

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient
 3. well-calibrated
 4. well-powered
- **Problem:** current methods fail to achieve all of these

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient
 3. well-calibrated
 4. well-powered
- **Problem:** current methods fail to achieve all of these
- This talk: **split predictive checks** (SPCs) *can be the answer!*

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient
 3. well-calibrated
 4. well-powered
- **Problem:** current methods fail to achieve all of these
- This talk: **split predictive checks** (SPCs) *can be the answer!*
 - ▶ develop **single SPCs** and **divided SPCs**

Model checking: can we have it all?

- **Model criticism** a crucial part of Bayesian model building:
 - ▶ Evaluates current model fit, insights for model elaboration, ...
- Want a model criticism method to be...
 1. general-purpose
 2. computationally efficient
 3. well-calibrated
 4. well-powered
- **Problem:** current methods fail to achieve all of these
- This talk: **split predictive checks** (SPCs) *can be the answer!*
 - ▶ develop **single SPCs** and **divided SPCs**
 - ▶ provide asymptotic theory showing good **calibration & power**

Bayesian model checking

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_0$.

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_\circ$
- Model $\{P_\theta : \theta \in \Theta\}$

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_*$
- Model $\{P_\theta : \theta \in \Theta\}$
- KL-optimal parameter $\theta_\star := \sup_{\theta \in \Theta} \mathbb{E}\{\log p_\theta(x_1)\}$

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_\circ$
- Model $\{P_\theta : \theta \in \Theta\}$
- KL-optimal parameter $\theta_\star := \sup_{\theta \in \Theta} \mathbb{E}\{\log p_\theta(x_1)\}$
- **Goal:** measure the mismatch between P_\circ and P_{θ_\star}

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_\circ$
- Model $\{P_\theta : \theta \in \Theta\}$
- KL-optimal parameter $\theta_\star := \sup_{\theta \in \Theta} \mathbb{E}\{\log p_\theta(x_1)\}$
- **Goal:** measure the mismatch between P_\circ and P_{θ_\star}
 - $H_0 : P_{\theta_\star} = P_\circ$ (well-specified)

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_*$
- Model $\{P_\theta : \theta \in \Theta\}$
- KL-optimal parameter $\theta_\star := \sup_{\theta \in \Theta} \mathbb{E}\{\log p_\theta(x_1)\}$
- **Goal:** measure the mismatch between P_* and P_{θ_\star}
 - $H_0 : P_{\theta_\star} = P_*$ (well-specified)
 - $H_1 : P_\theta \neq P_* \forall \theta \in \Theta$ (misspecified)

Bayesian model checking

- Data $\mathbf{x} = (x_1, \dots, x_N)$ where $x_n \stackrel{\text{i.i.d.}}{\sim} P_0$
- Model $\{P_\theta : \theta \in \Theta\}$
- KL-optimal parameter $\theta_\star := \sup_{\theta \in \Theta} \mathbb{E}\{\log p_\theta(x_1)\}$
- **Goal:** measure the mismatch between P_0 and P_{θ_\star}
 - $H_0 : P_{\theta_\star} = P_0$ (well-specified)
 - $H_1 : P_\theta \neq P_0 \ \forall \theta \in \Theta$ (misspecified)
- p-values provide measure of fit only if they are well-calibrated!

Predictive checks


Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$

Predictive checks


- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


*user-specified
statistic*

Predictive checks

- A predictive p-value is given by


$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


*user-specified
statistic*

- Standard choices of m

Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$



*user-specified
statistic*

- Standard choices of m

- ▶ Prior PCs: $m_{\text{prior}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi_0(d\theta)$ [Box (1980)]

Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


*user-specified
statistic*


- Standard choices of m

- ▶ Prior PCs: $m_{\text{prior}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi_0(d\theta)$ [Box (1980)]

- ▶ Posterior PCs: $m_{\text{post}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi(d\theta | \mathbf{x})$ [Rubin (1984)]

Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


user-specified statistic

- Standard choices of m

- ▶ Prior PCs: $m_{\text{prior}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi_0(d\theta)$ [Box (1980)]


- ▶ Posterior PCs: $m_{\text{post}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi(d\theta | \mathbf{x})$ [Rubin (1984)]

- Posterior predictive checks (PPCs) p-value is defined as

$$p_{PPC}(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}(\cdot | \mathbf{x})} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$

Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


user-specified statistic

- Standard choices of m

- ▶ Prior PCs: $m_{\text{prior}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi_0(d\theta)$ [Box (1980)]

- ▶ Posterior PCs: $m_{\text{post}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi(d\theta | \mathbf{x})$ [Rubin (1984)]

- Posterior predictive checks (PPCs) p-value is defined as

$$p_{PPC}(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}(\cdot | \mathbf{x})} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


- ✓ Flexible, general-purpose, and computationally efficient

Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$

user-specified statistic



- Standard choices of m

- ▶ Prior PCs: $m_{\text{prior}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi_0(d\theta)$ [Box (1980)]
- ▶ Posterior PCs: $m_{\text{post}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi(d\theta | \mathbf{x})$ [Rubin (1984)]

- Posterior predictive checks (PPCs) p-value is defined as

$$p_{PPC}(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}(\cdot | \mathbf{x})} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$


- ✓ Flexible, general-purpose, and computationally efficient
- ✗ not well-calibrated, conservative p-values [Robins et al (2000), Bayarri & Berger (2000), Moran et al (2022)]

Predictive checks

- A predictive p-value is given by

$$p(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$

*user-specified
statistic*



- Standard choices of m


► Prior PCs: $m_{\text{prior}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi_0(d\theta)$ [Box (1980)]

► Posterior PCs: $m_{\text{post}}(\mathbf{x}_{\text{pred}}) := \int p_{\theta}(\mathbf{x}_{\text{pred}}) \Pi(d\theta | \mathbf{x})$ [Rubin (1984)]

- Posterior predictive checks (PPCs) p-value is defined as

$$p_{PPC}(\mathbf{x}) = \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}(\cdot | \mathbf{x})} \{ T(\mathbf{x}) > T(\mathbf{x}_{\text{pred}}) \}$$

*double use
of data*



- ✓ Flexible, general-purpose, and computationally efficient

- ✗ not well-calibrated, conservative p-values [Robins et al (2000), Bayarri & Berger (2000), Moran et al (2022)]

PPCs are not well-calibrated

PPCs are not well-calibrated

- Model: $P_\theta = \text{Poiss}(\theta)$

PPCs are not well-calibrated

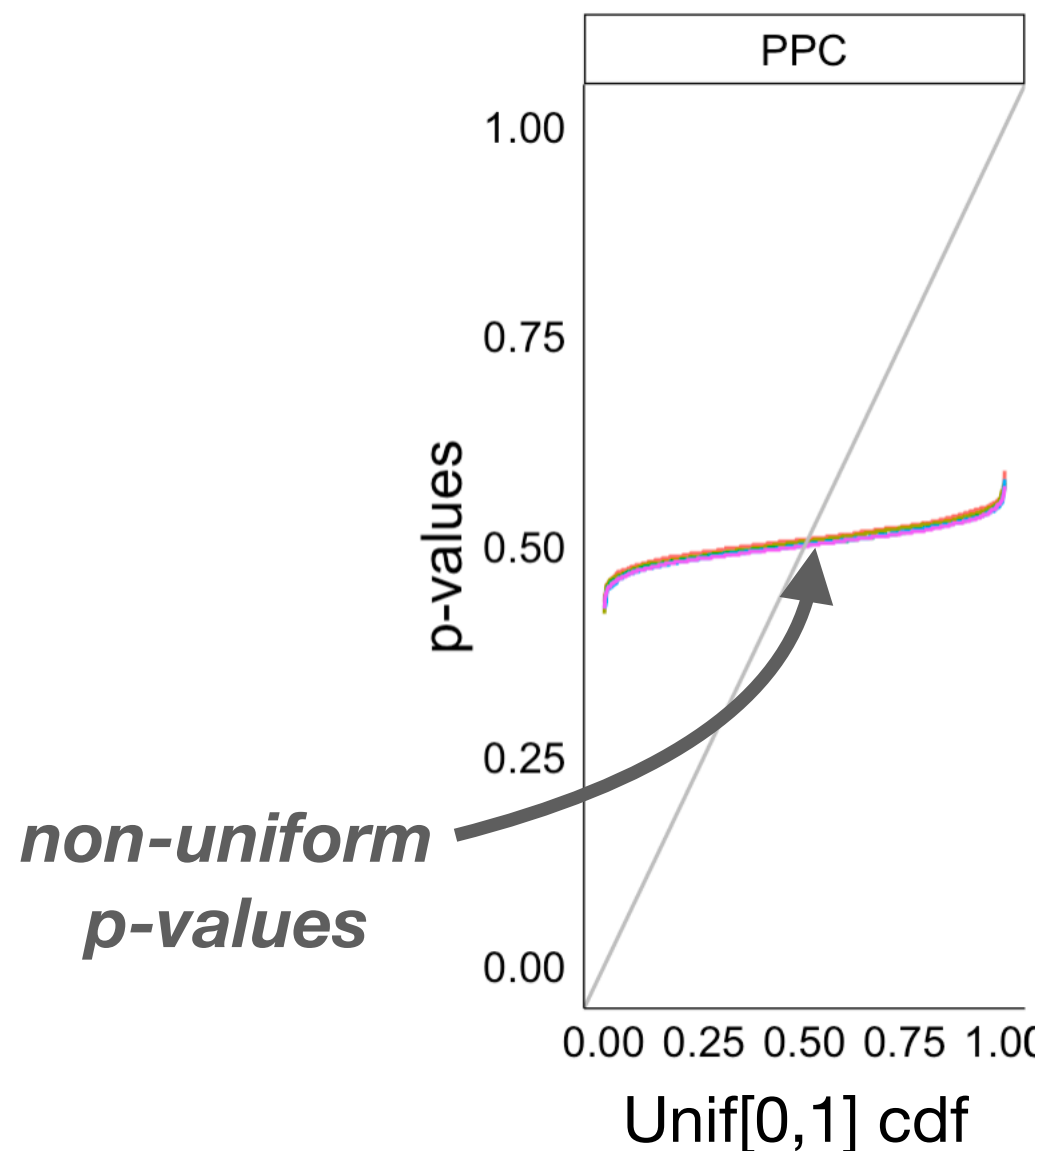
- Model: $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{Poiss}(2)$

PPCs are not well-calibrated

- Model: $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{Poiss}(2)$
- So, model is well-specified [H_0 holds]

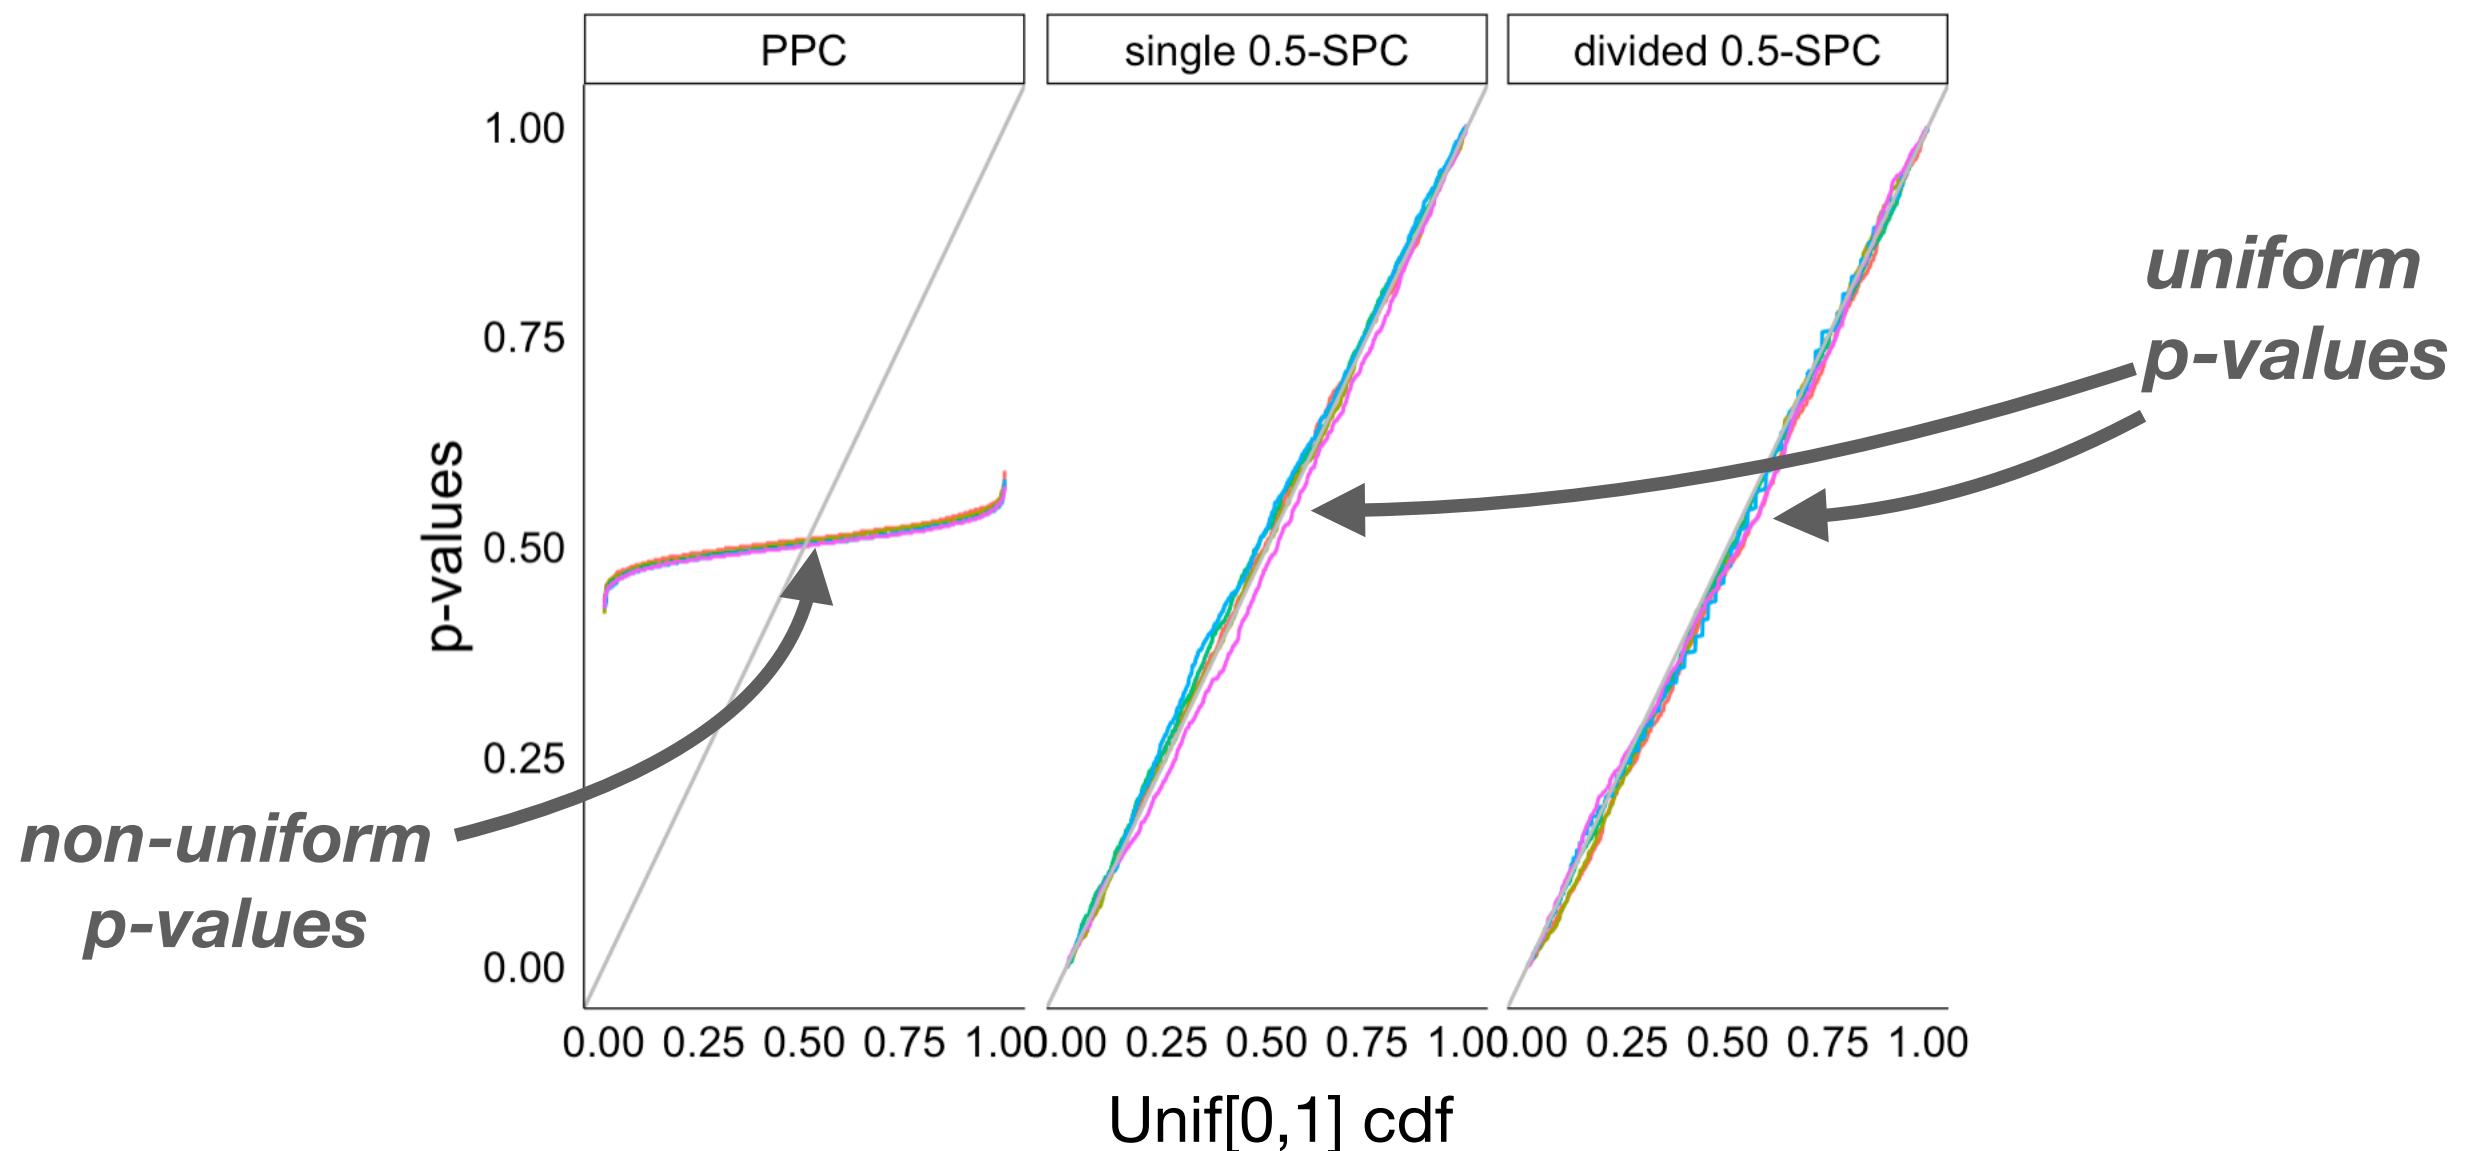
PPCs are not well-calibrated

- Model: $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{Poiss}(2)$
- So, model is well-specified [H_0 holds]



PPCs are not well-calibrated

- Model: $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{Poiss}(2)$
- So, model is well-specified [H_0 holds]



PPCs produce conservative p-values

PPCs produce conservative p-values

- Geometric model for airline delay data for 2013 NYC departures

$$y \sim \text{Geom}(\theta), \quad \theta \sim \text{Beta}(0.1, 0.2)$$

PPCs produce conservative p-values

- Geometric model for airline delay data for 2013 NYC departures

$$y \sim \text{Geom}(\theta), \quad \theta \sim \text{Beta}(0.1, 0.2)$$

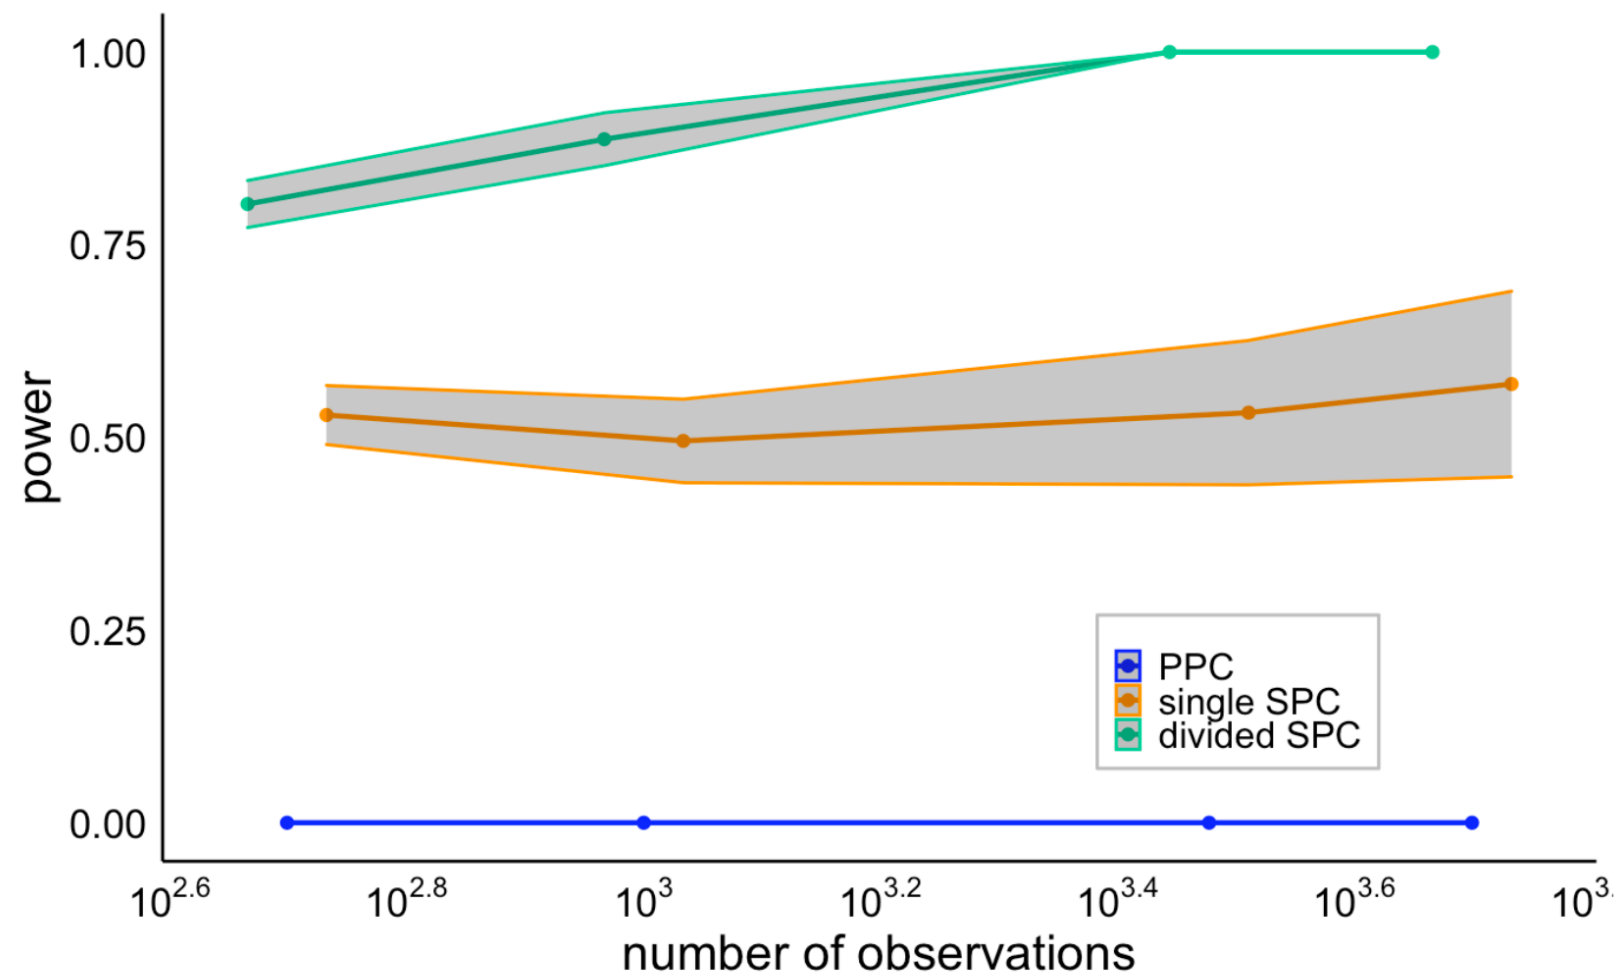
- Model is misspecified [H_1 holds]

PPCs produce conservative p-values

- Geometric model for airline delay data for 2013 NYC departures

$$y \sim \text{Geom}(\theta), \quad \theta \sim \text{Beta}(0.1, 0.2)$$

- Model is misspecified [H_1 holds]

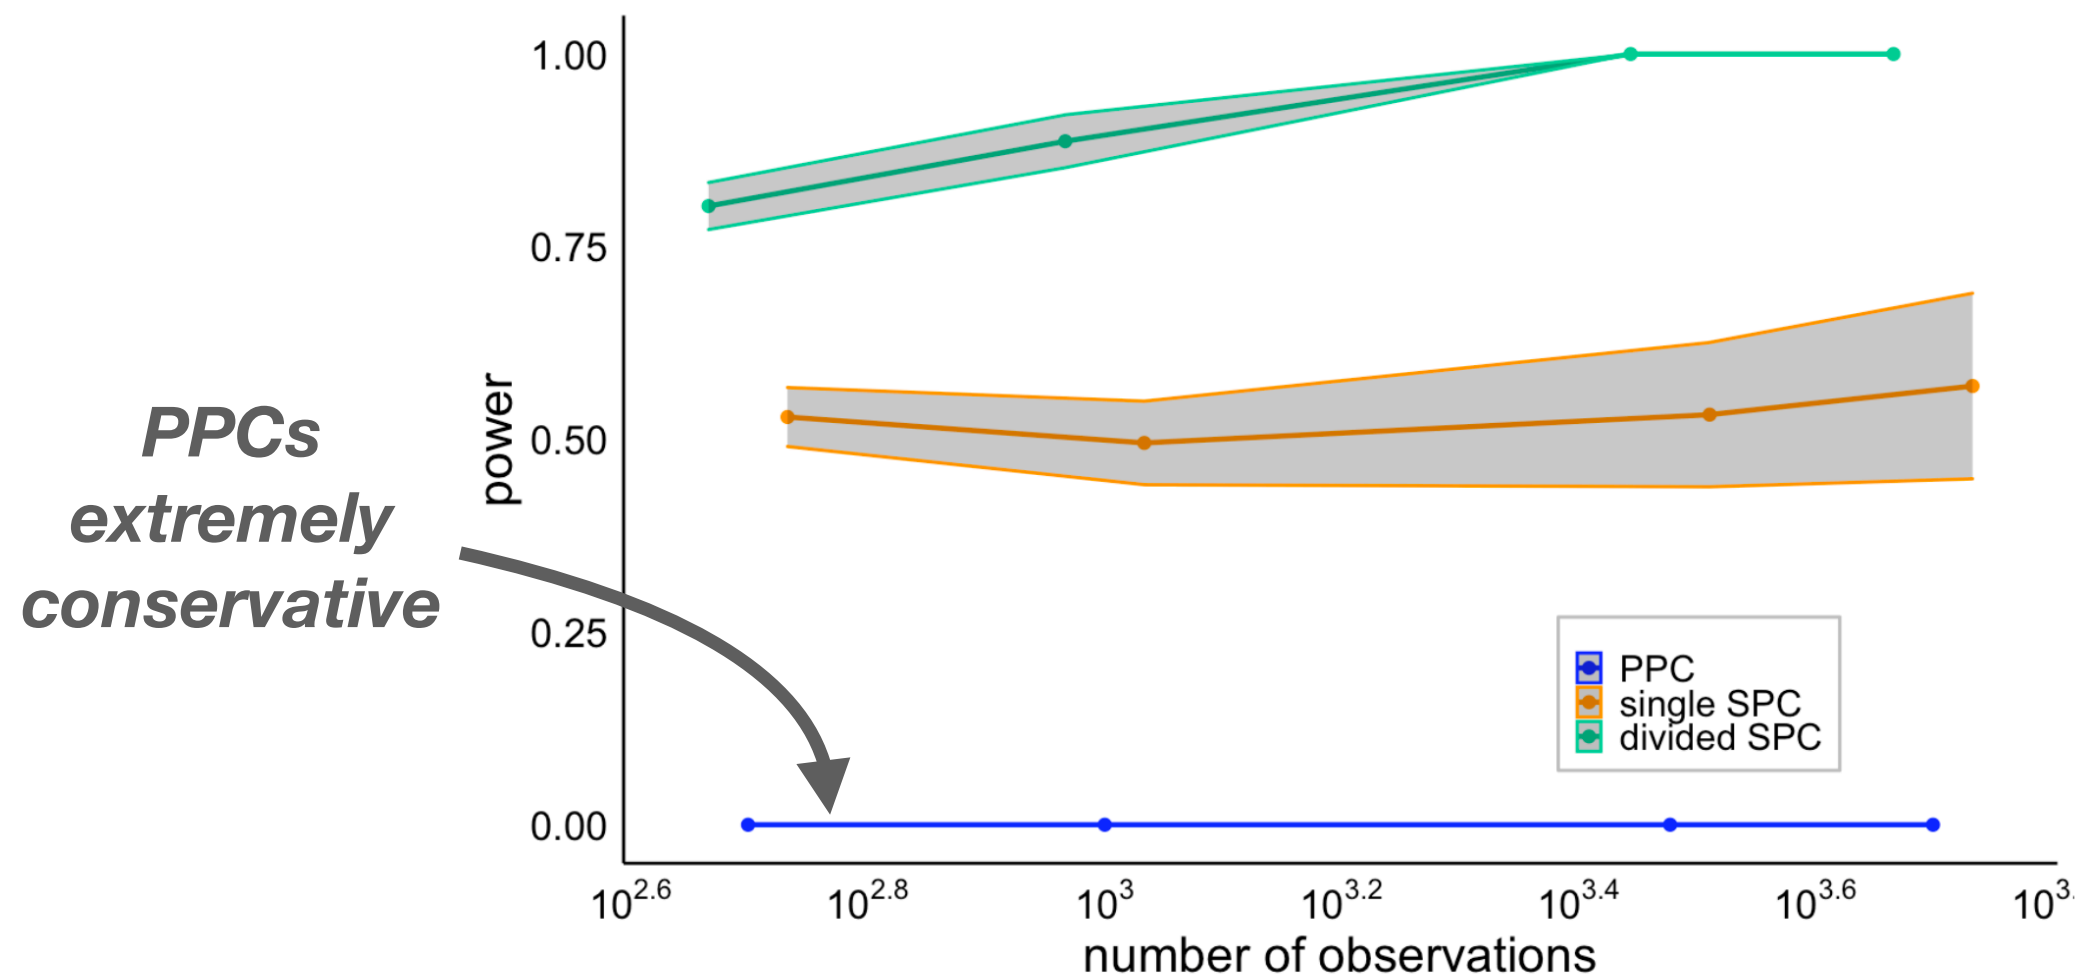


PPCs produce conservative p-values

- Geometric model for airline delay data for 2013 NYC departures

$$y \sim \text{Geom}(\theta), \quad \theta \sim \text{Beta}(0.1, 0.2)$$

- Model is misspecified [H_1 holds]

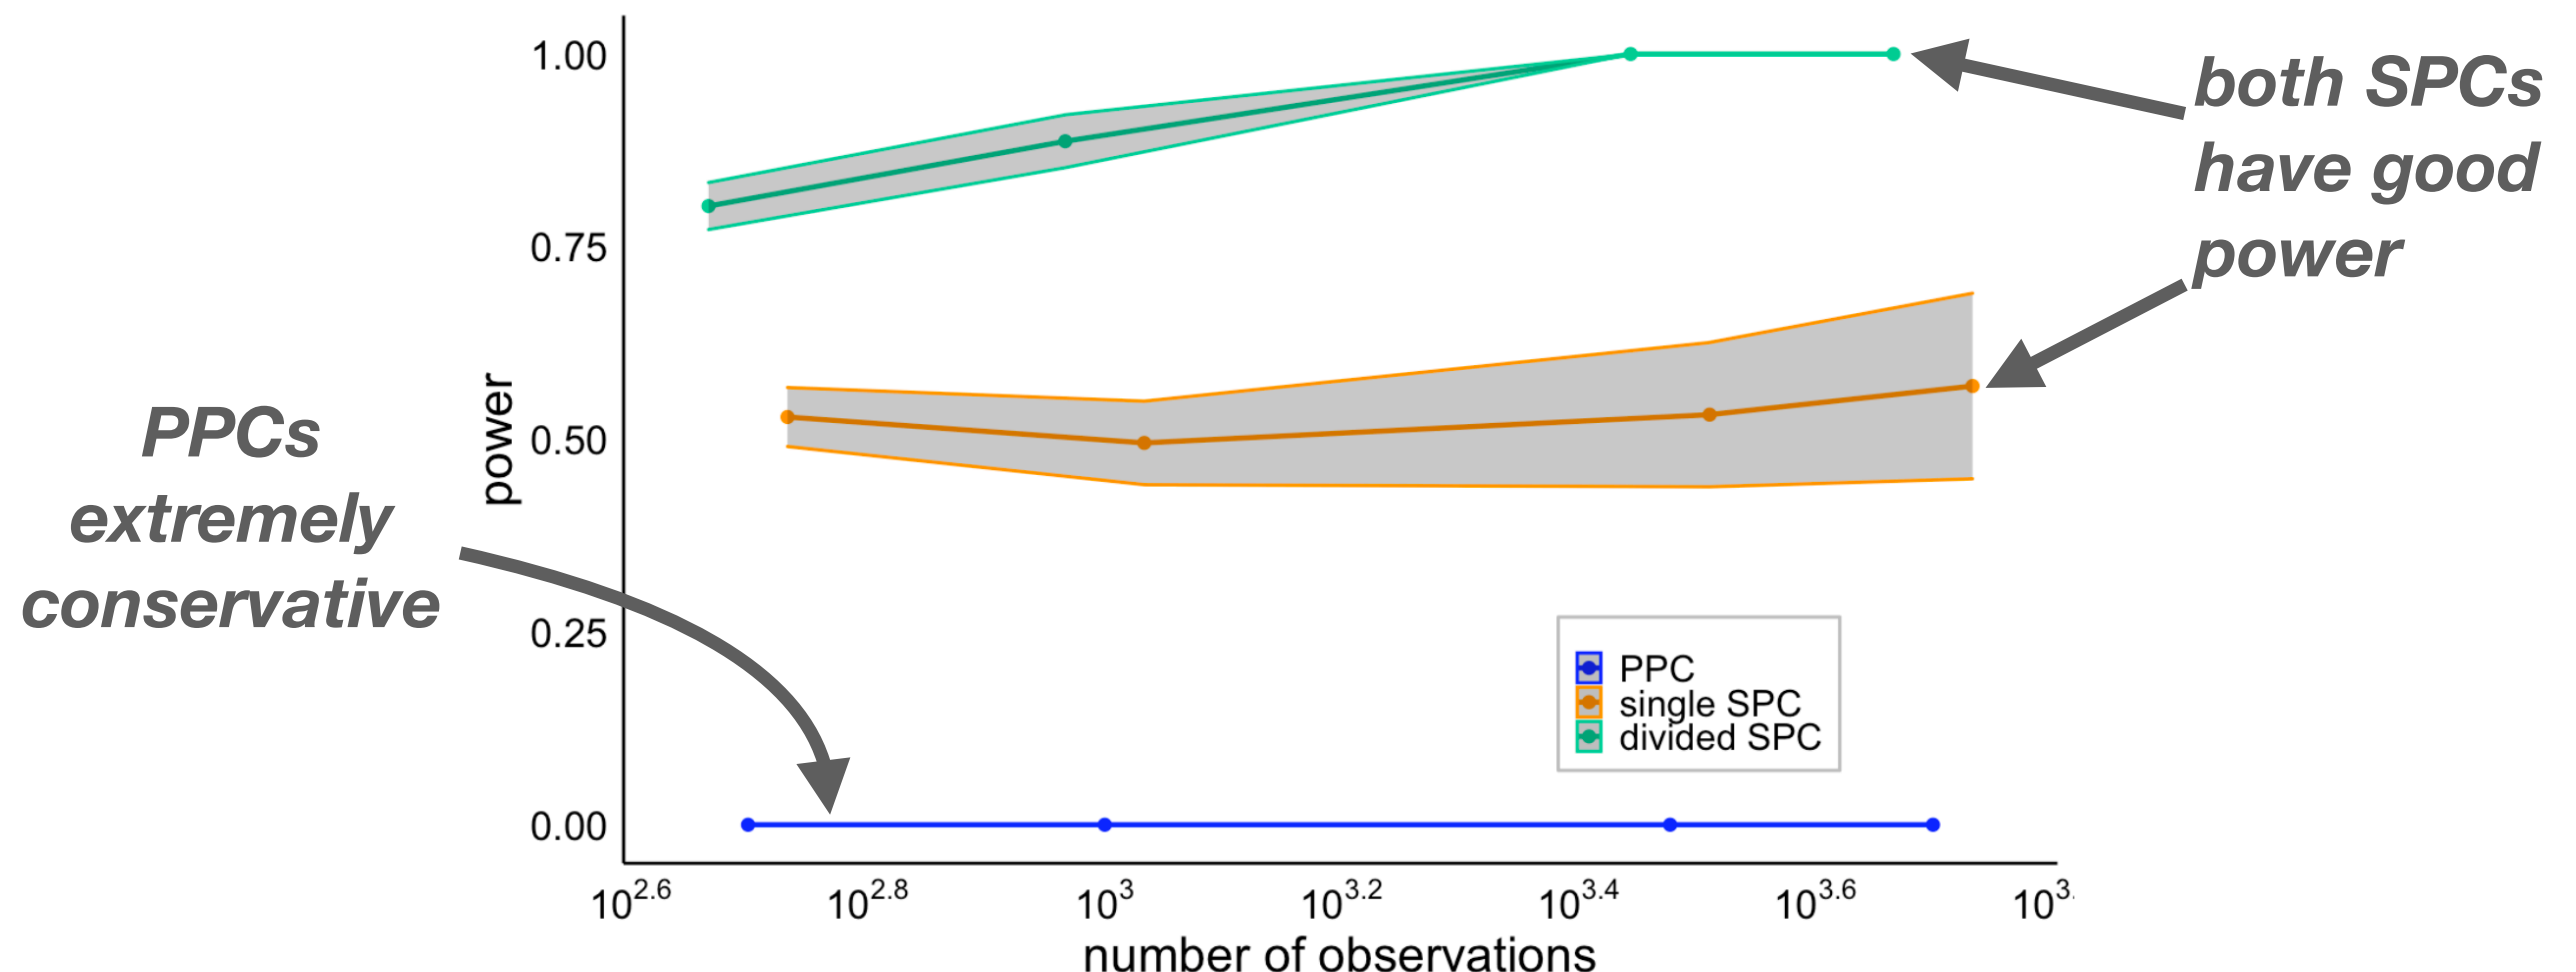


PPCs produce conservative p-values

- Geometric model for airline delay data for 2013 NYC departures

$$y \sim \text{Geom}(\theta), \quad \theta \sim \text{Beta}(0.1, 0.2)$$

- Model is misspecified [H_1 holds]



Single split predictive checks

Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data


Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data

X


Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data



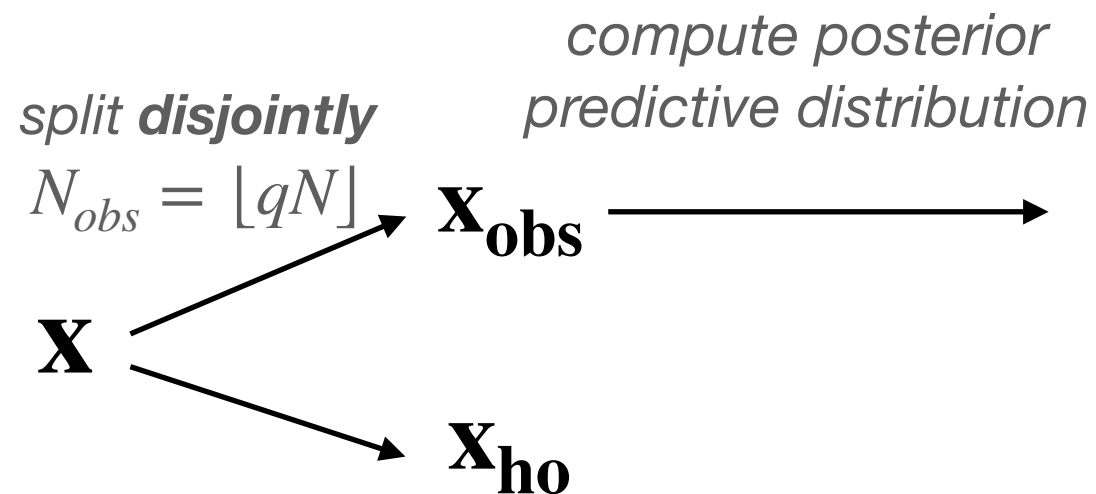
Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data



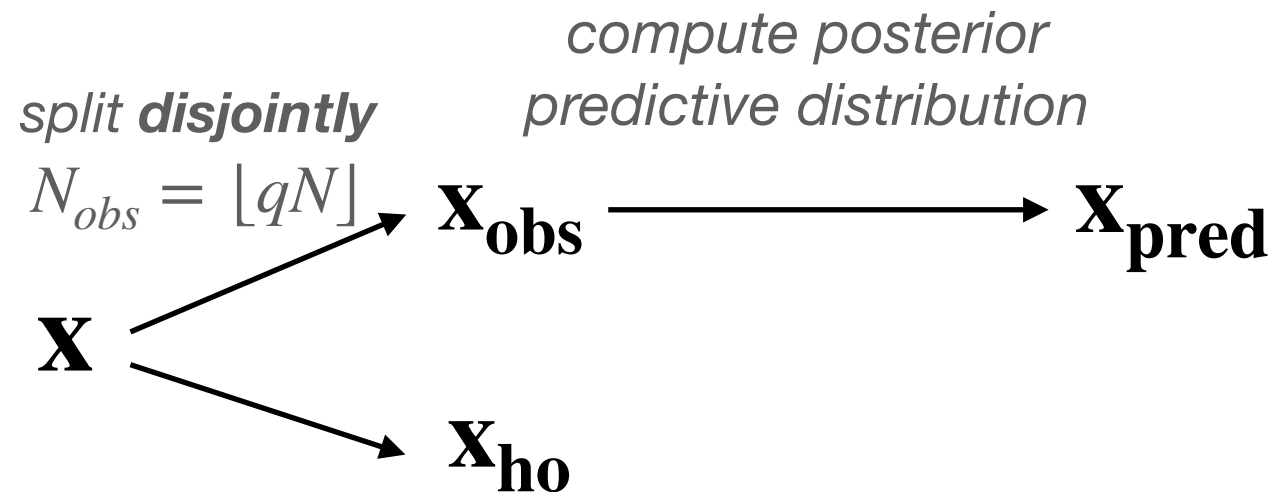
Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data




Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data




Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data



Single split predictive checks


- Single SPCs avoid double use of data by splitting data into “train” and “test” data



- Single q-SPC p-value [see also Moran et al (2022)]:

Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data




- Single q-SPC p-value [see also Moran et al (2022)]:

$$p_{\text{SPC}}(\mathbf{X}) := \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}^{\text{SPC}}(\cdot | \mathbf{x}_{\text{obs}})} \{ T(\mathbf{x}_{\text{ho}}) > T(\mathbf{x}_{\text{pred}}) \}$$

Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data



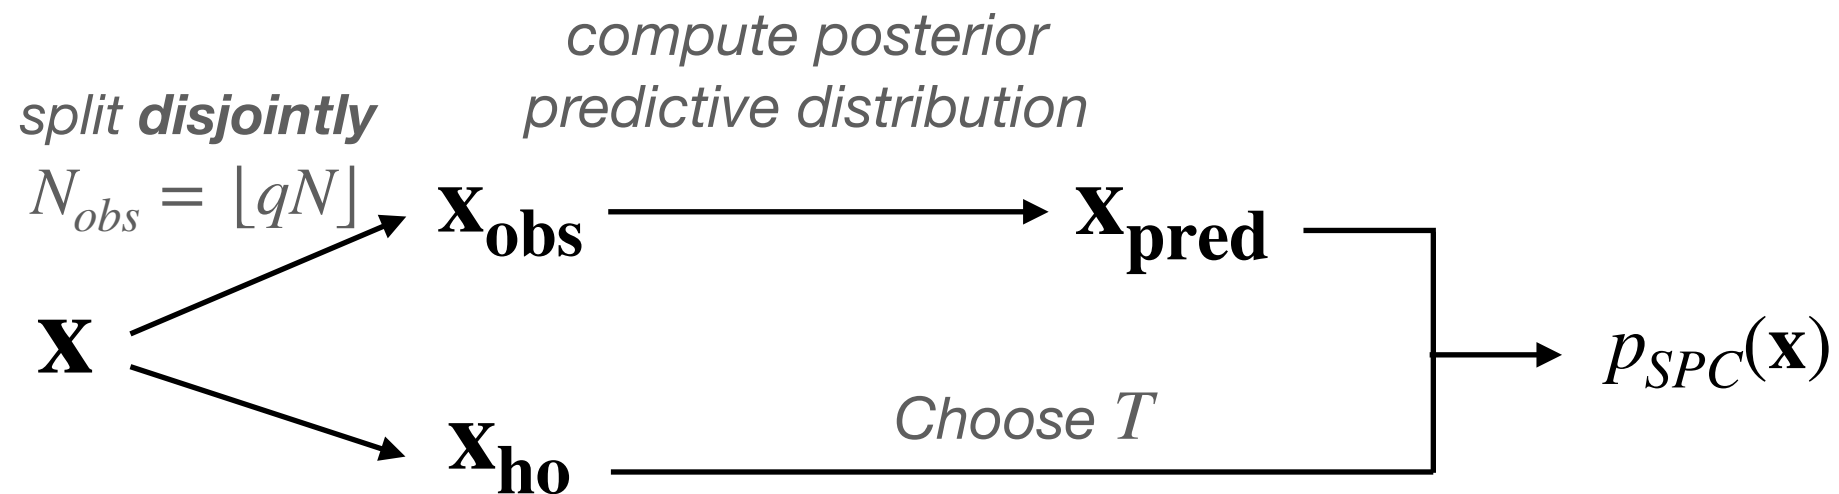
- Single q-SPC p-value [see also Moran et al (2022)]:

$$p_{\text{SPC}}(\mathbf{x}) := \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}^{\text{SPC}}(\cdot | \mathbf{x}_{\text{obs}})} \{ T(\mathbf{x}_{\text{ho}}) > T(\mathbf{x}_{\text{pred}}) \}$$

- ✓ flexible, computationally efficient, and general-purpose

Single split predictive checks

- Single SPCs avoid double use of data by splitting data into “train” and “test” data



- Single q-SPC p-value [see also Moran et al (2022)]:

$$p_{\text{SPC}}(\mathbf{x}) := \mathbb{P}_{\mathbf{x}_{\text{pred}} \sim m_{\text{post}}^{\text{SPC}}(\cdot | \mathbf{x}_{\text{obs}})} \{ T(\mathbf{x}_{\text{ho}}) > T(\mathbf{x}_{\text{pred}}) \}$$

- ✓ flexible, computationally efficient, and general-purpose

Question 1: How is calibration and power?

Asymptotic power and calibration of single SPCs

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Theorem [LH22]: if $\nu_{\circ} = \nu(\theta_{\star})$, then $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Theorem [LH22]: if $\nu_{\circ} = \nu(\theta_{\star})$, then $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

ρ depends on q , Fisher information matrices and asymptotic variances of T

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Theorem [LH22]: if $\nu_{\circ} = \nu(\theta_{\star})$, then $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Theorem [LH22]: if $\nu_{\circ} = \nu(\theta_{\star})$, then $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$
- if model misspecified, then in general $\rho \neq 1$:

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Theorem [LH22]: if $\nu_{\circ} = \nu(\theta_{\star})$, then $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$
- if model misspecified, then in general $\rho \neq 1$:
 - when uncertainty is under-estimated, $\rho \uparrow \infty$ and power $\uparrow 1$

Asymptotic power and calibration of single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$, then power $\xrightarrow{P} 1$

Theorem [LH22]: if $\nu_{\circ} = \nu(\theta_{\star})$, then $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$
- if model misspecified, then in general $\rho \neq 1$:
 - when uncertainty is under-estimated, $\rho \uparrow \infty$ and power $\uparrow 1$
 - when uncertainty is over-estimated, $\rho \downarrow 0$ and power $\downarrow 0$

**Single SPCs have asymptotic power 1
when misspecification is significant**

Single SPCs have asymptotic power 1 when misspecification is significant

- Model $P_\theta = \text{Poiss}(\theta)$

Single SPCs have asymptotic power 1 when misspecification is significant

- Model $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{NegBin}(2, \tau)$ and $P_\circ = \text{Binom}(30, p)$

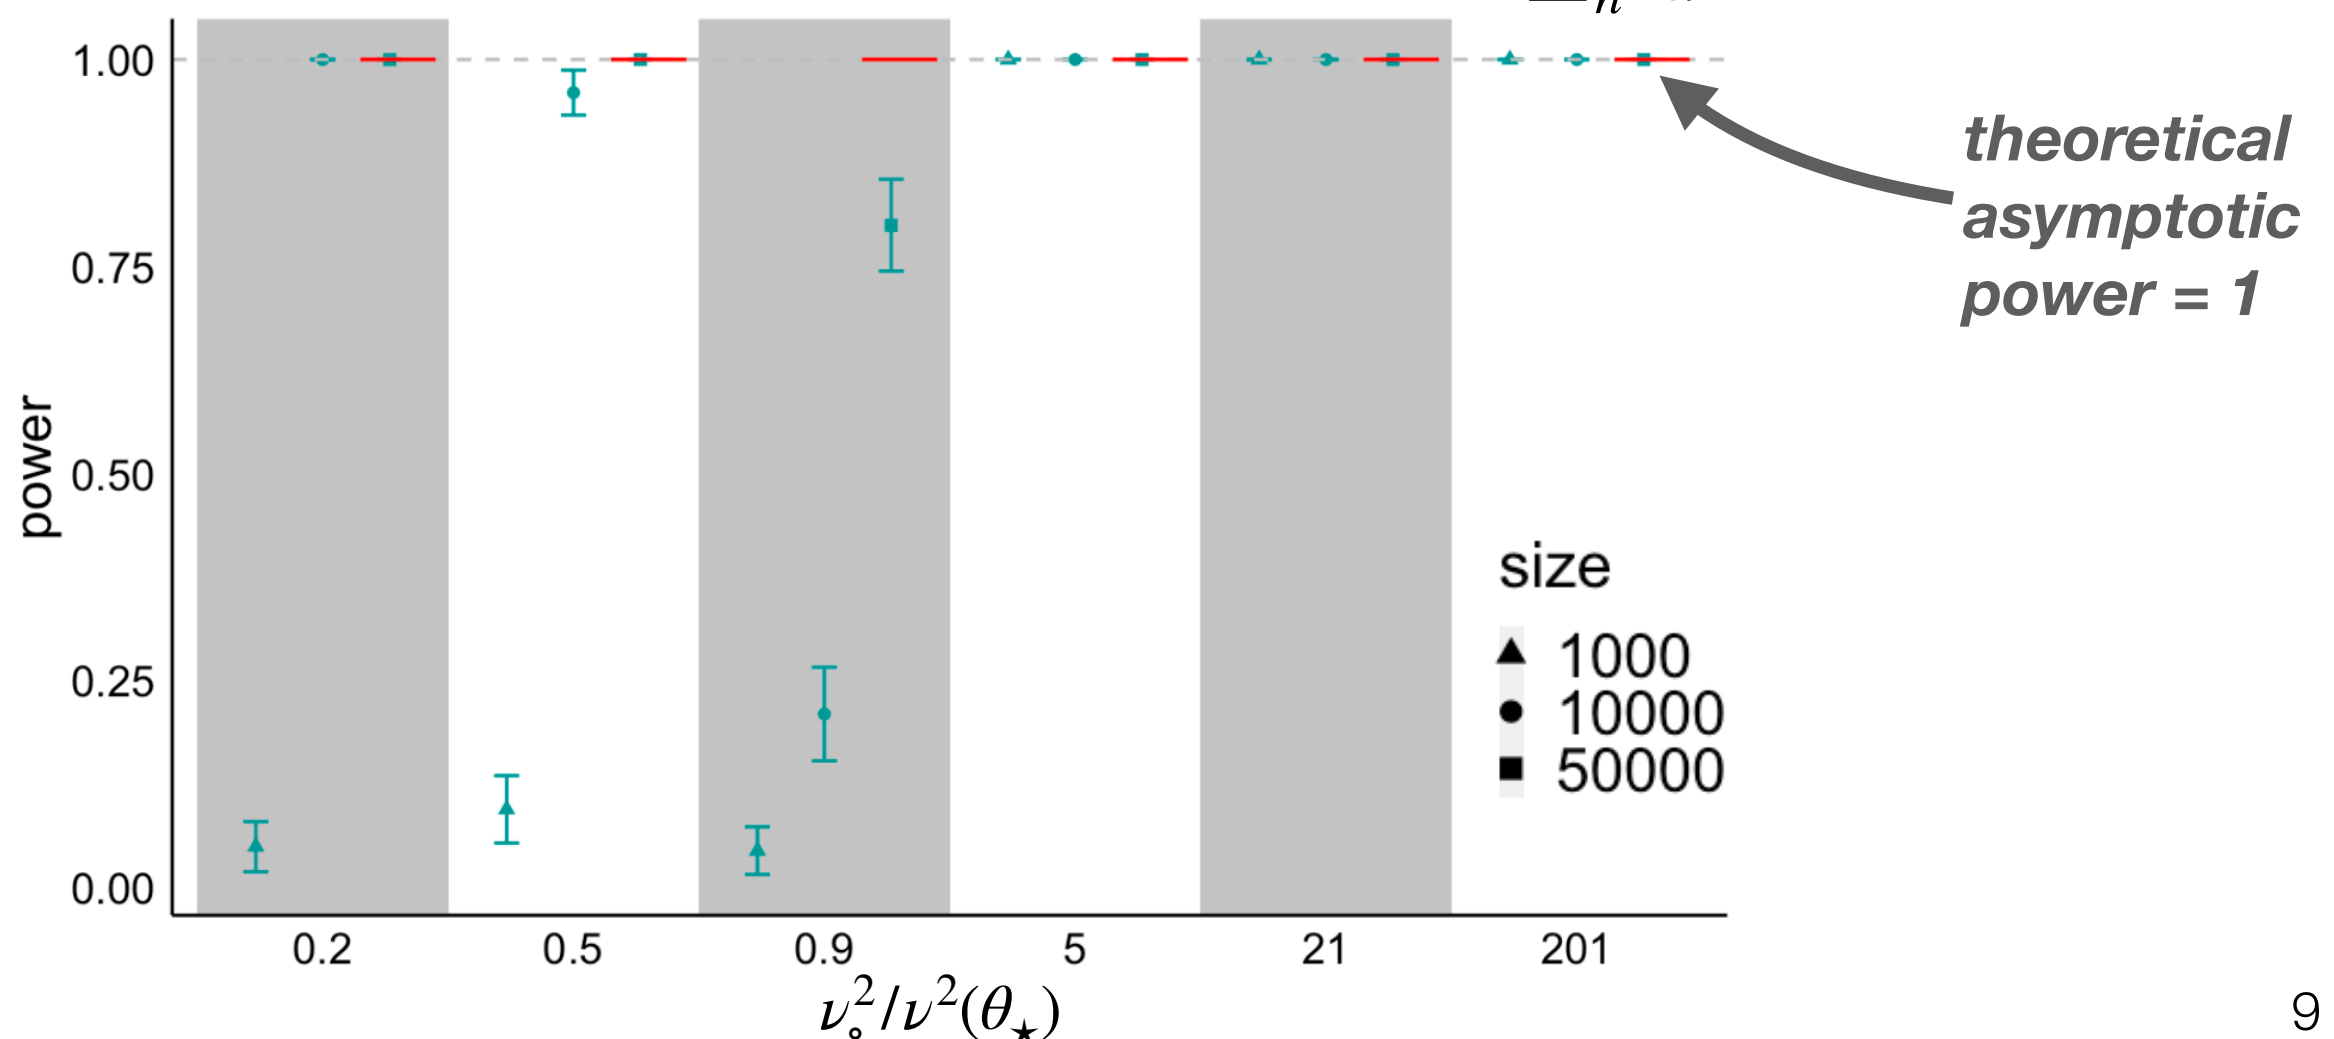
Single SPCs have asymptotic power 1 when misspecification is significant

- Model $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{NegBin}(2, \tau)$ and $P_\circ = \text{Binom}(30, p)$
- Recall: when $\nu_\circ / \nu(\theta_\star) \neq 1$, asymptotic power = 1

Single SPCs have asymptotic power 1 when misspecification is significant

- Model $P_\theta = \text{Poiss}(\theta)$
- Generate data from $P_\circ = \text{NegBin}(2, \tau)$ and $P_\circ = \text{Binom}(30, p)$
- Recall: when $\nu_\circ/\nu(\theta_\star) \neq 1$, asymptotic power = 1

Single 0.5-SPC, second moment statistic $T(\mathbf{x}) = \sum_n x_n^2/N$



Possible poor power of single SPCs when misspecification is subtle

Possible poor power of single SPCs when misspecification is subtle

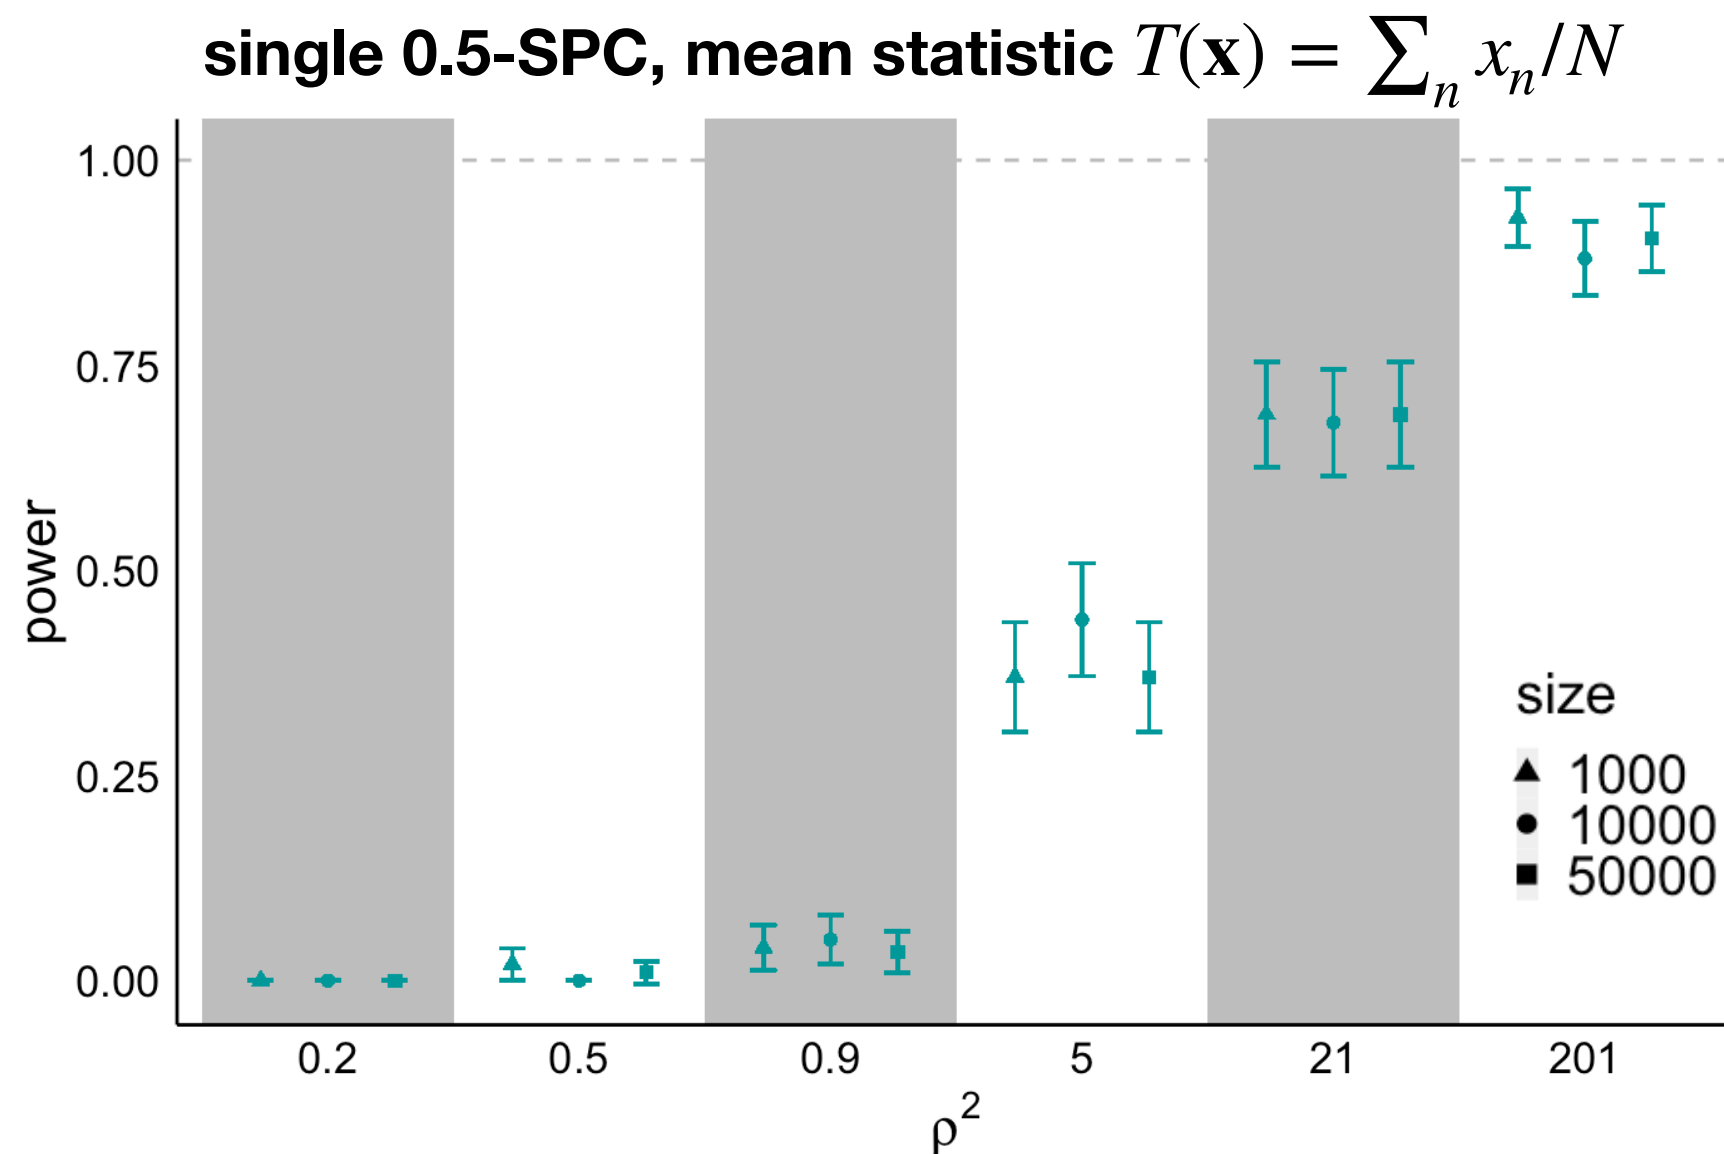
- Recall: if $\nu_{\circ} = \nu(\theta_{\star})$, asymptotic power depends on ρ

Possible poor power of single SPCs when misspecification is subtle

- Recall: if $\nu_{\circ} = \nu(\theta_{\star})$, asymptotic power depends on ρ
- For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, have $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$

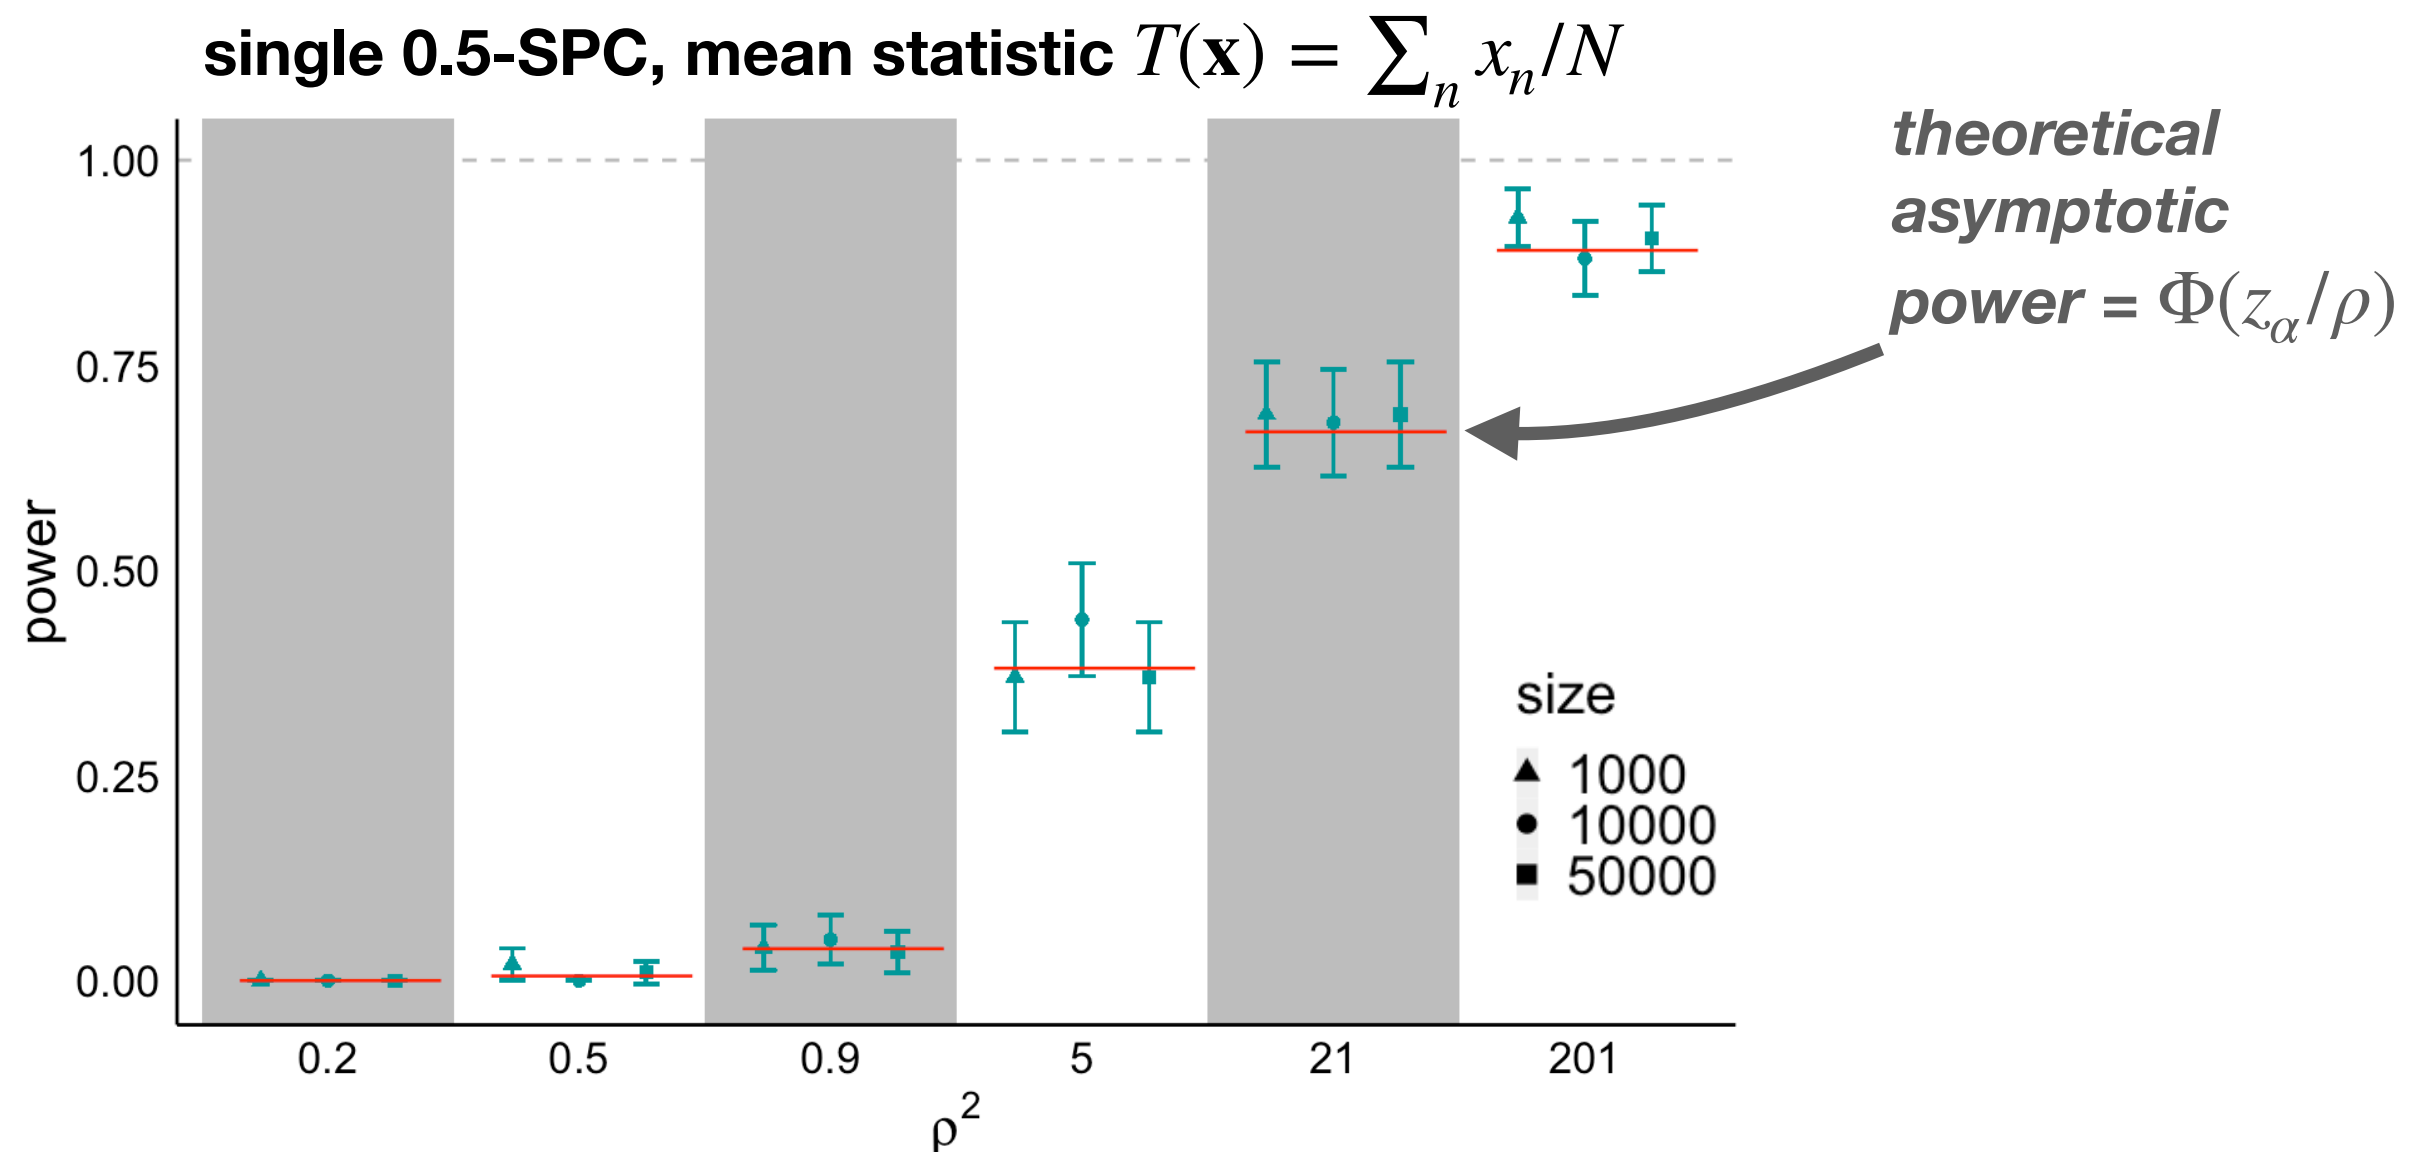
Possible poor power of single SPCs when misspecification is subtle

- Recall: if $\nu_{\circ} = \nu(\theta_{\star})$, asymptotic power depends on ρ
- For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, have $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$



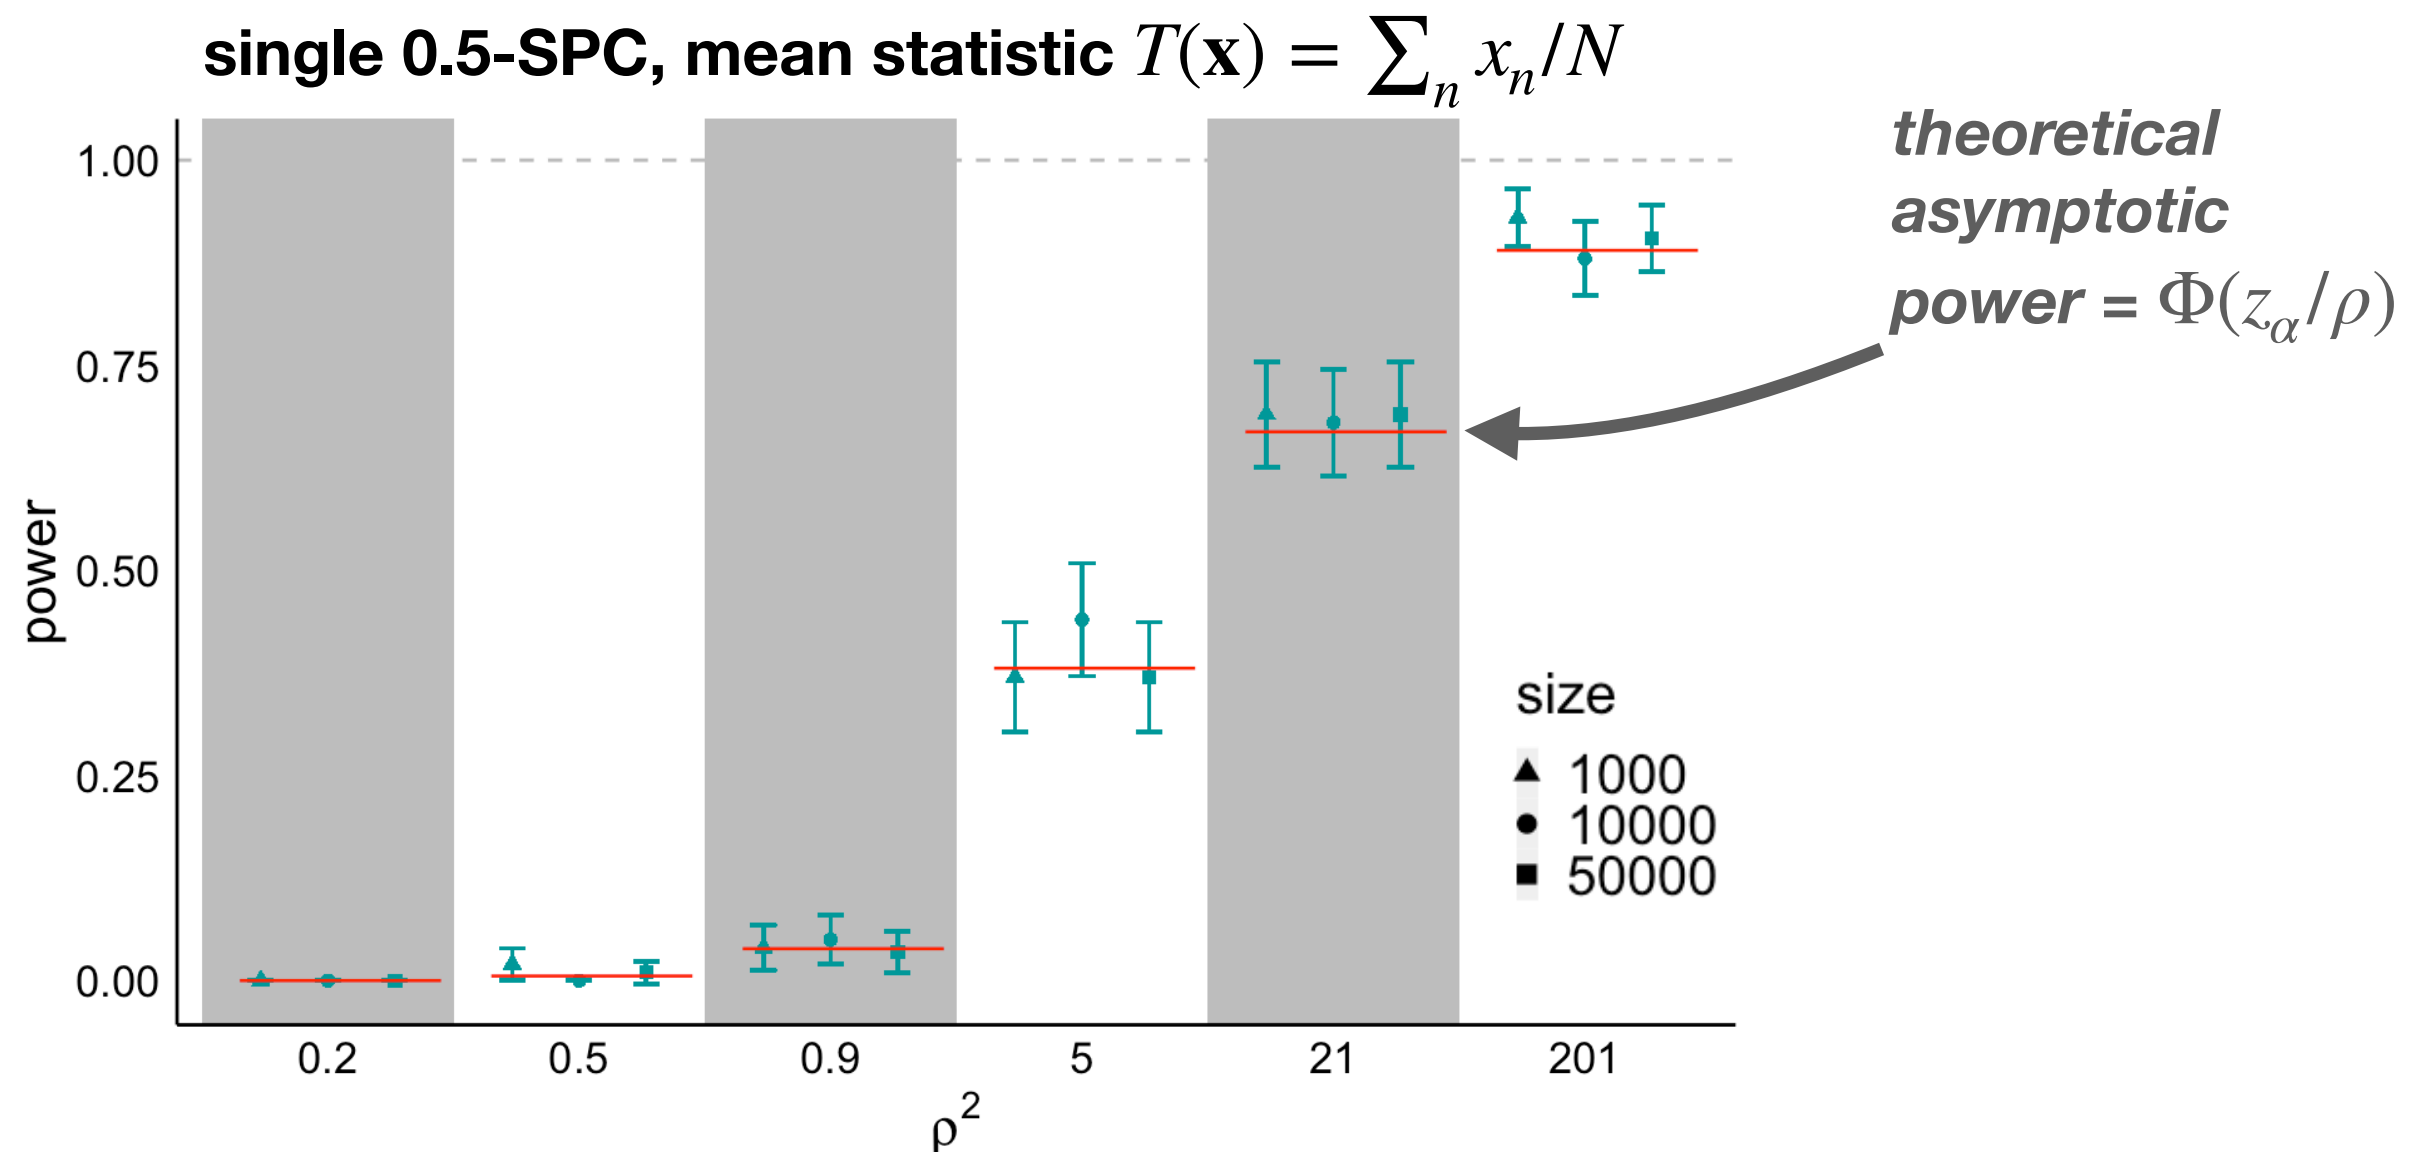
Possible poor power of single SPCs when misspecification is subtle

- Recall: if $\nu_\circ = \nu(\theta_\star)$, asymptotic power depends on ρ
- For mean statistic $T(\mathbf{x}) = \sum_n x_n/N$, have $\nu_\circ = \nu(\theta_\star)$ and $\rho = \sigma_\circ/\sigma(\theta_\star)$



Possible poor power of single SPCs when misspecification is subtle

- Recall: if $\nu_{\circ} = \nu(\theta_{\star})$, asymptotic power depends on ρ
- For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, have $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$



Question 2: Can we improve power to detect subtle misspecification?

An SPC with better power

An SPC with better power

- **Goal:** propose an SPC that yields larger power under ***subtle misspecification*** – while maintaining calibration properties

An SPC with better power

- **Goal:** propose an SPC that yields larger power under ***subtle misspecification*** – while maintaining calibration properties
- From previous theorems, we expect:

An SPC with better power

- **Goal:** propose an SPC that yields larger power under ***subtle misspecification*** – while maintaining calibration properties
- From previous theorems, we expect:
 - when model is **misspecified**, single SPC p-values are **non-uniform** (also true when ρ is small)

An SPC with better power

- **Goal:** propose an SPC that yields larger power under ***subtle misspecification*** – while maintaining calibration properties
- From previous theorems, we expect:
 - ▶ when model is **misspecified**, single SPC p-values are **non-uniform** (also true when ρ is small)
 - ▶ when model is **correctly-specified**, single SPC p-values are **uniform**

An SPC with better power


- **Goal:** propose an SPC that yields larger power under ***subtle misspecification*** – while maintaining calibration properties
- From previous theorems, we expect:
 - ▶ when model is **misspecified**, single SPC p-values are **non-uniform** (also true when ρ is small)
 - ▶ when model is **correctly-specified**, single SPC p-values are **uniform**
- **Idea:** Test for ***uniformity*** of single SPC p-values

Divided SPCs


Divided SPCs

X

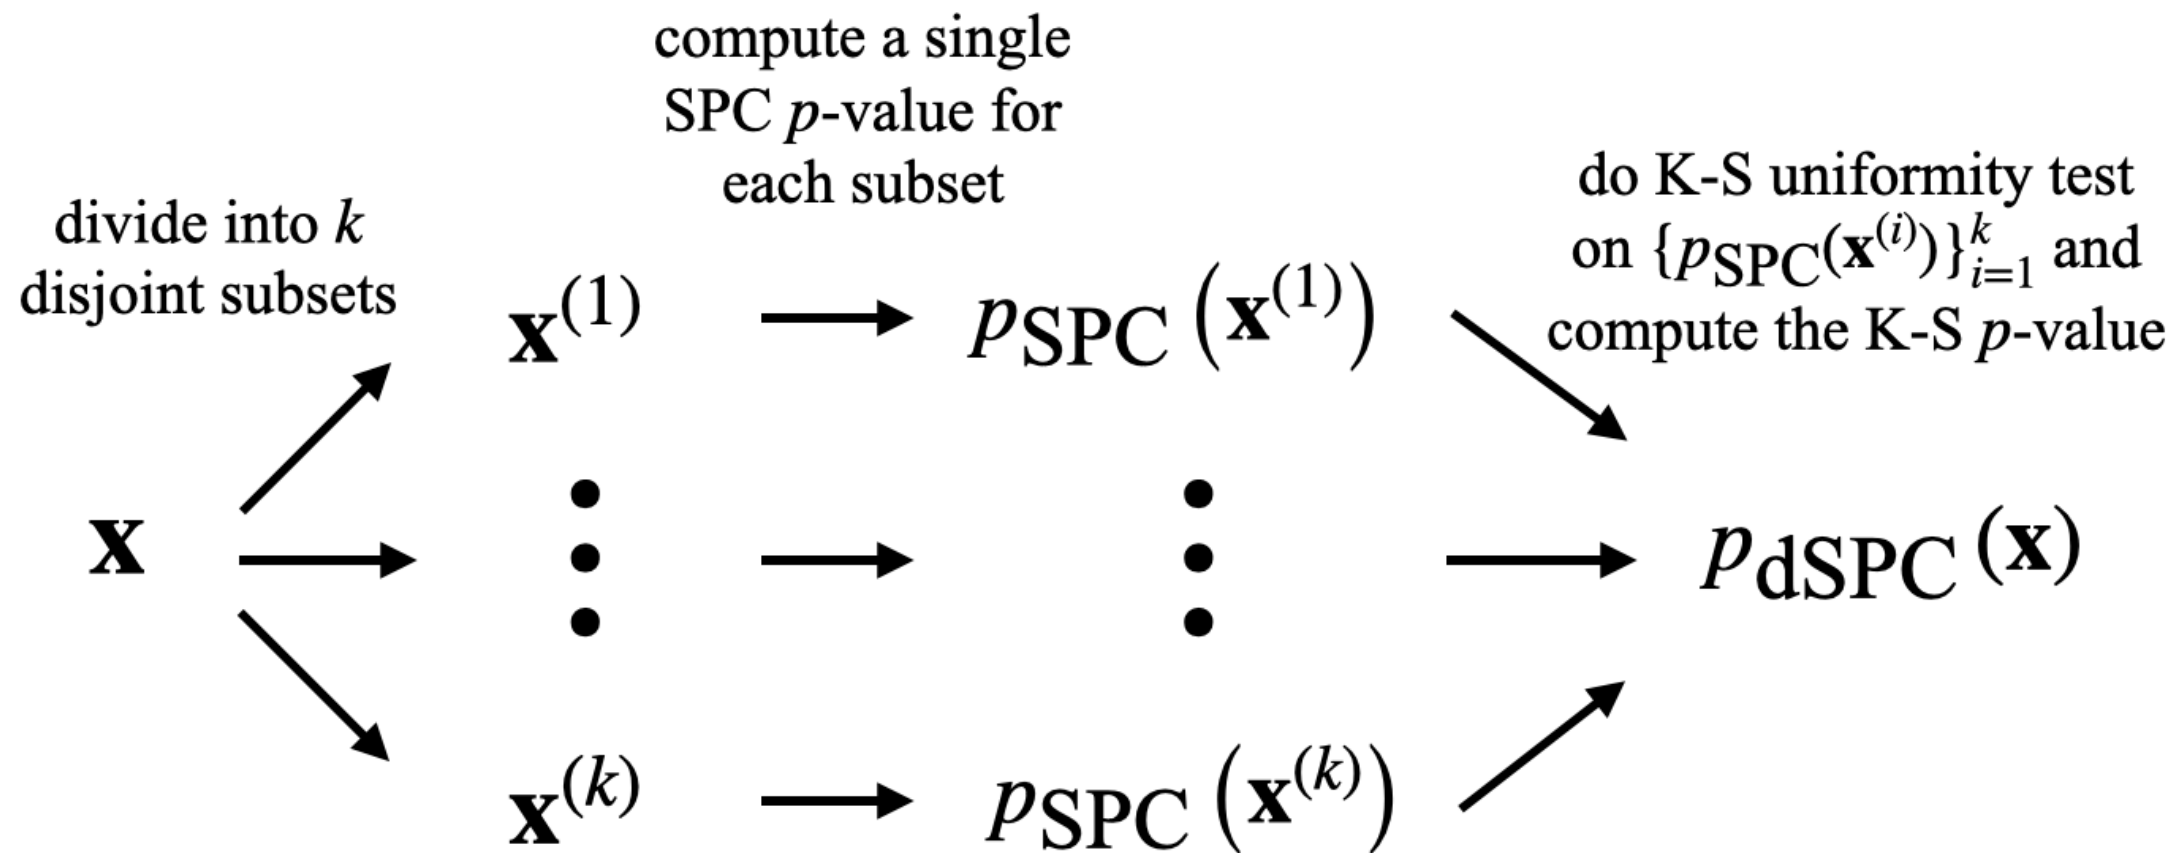
Divided SPCs



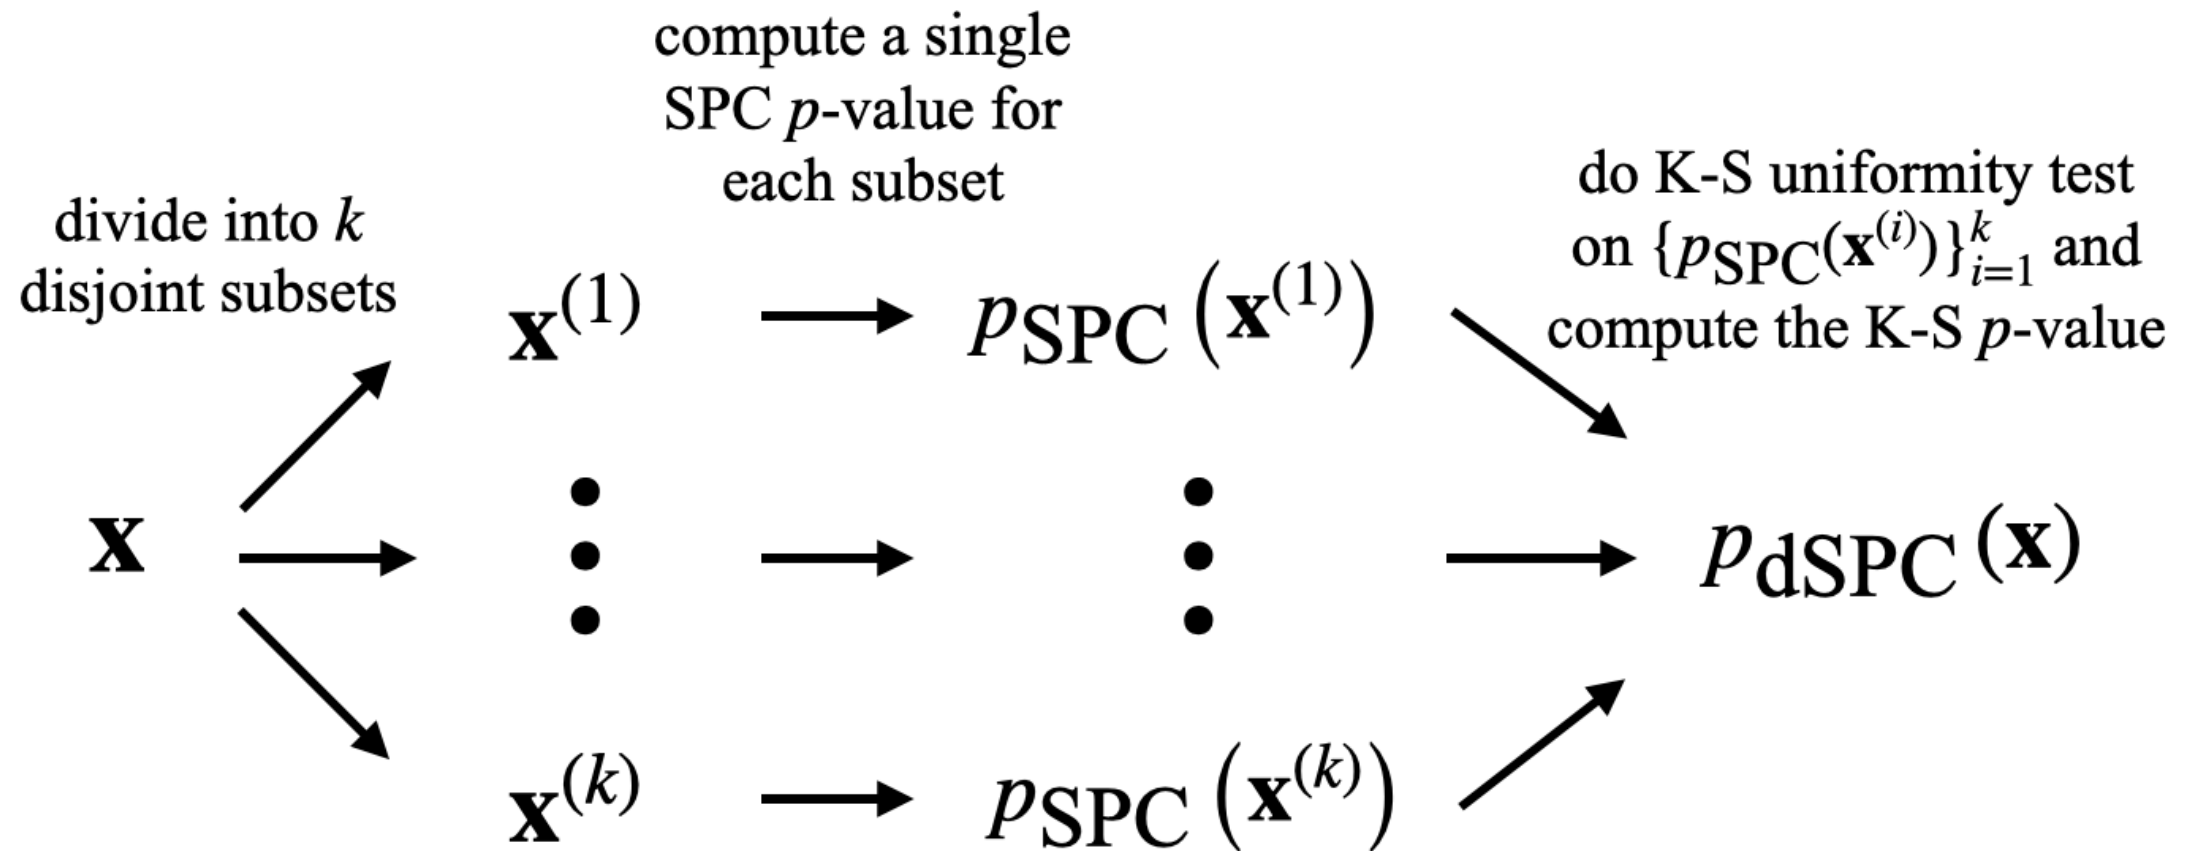
Divided SPCs



Divided SPCs



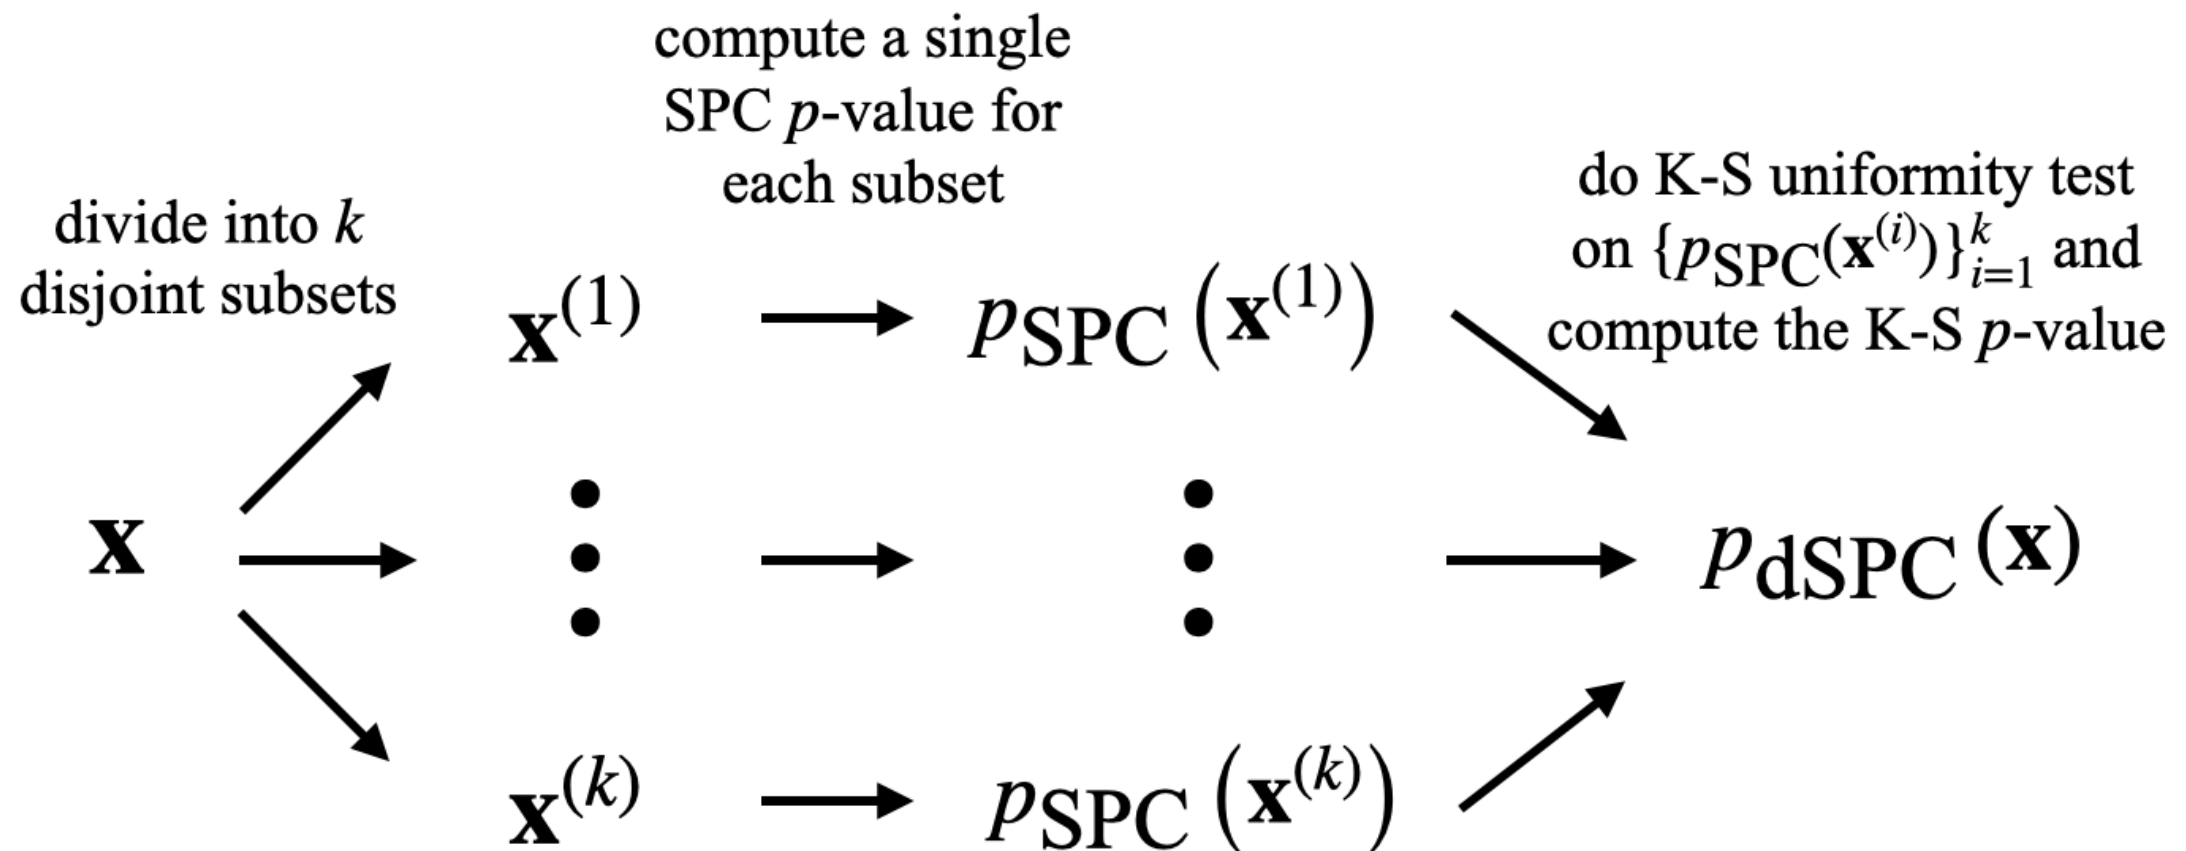
Divided SPCs



- Divided SPC p -value:

$$p_{\text{dSPC}}(\mathbf{X}) := p_{\text{KS}}(p_{\text{SPC}}(\mathbf{x}^{(1)}), \dots, p_{\text{SPC}}(\mathbf{x}^{(k)}))$$

Divided SPCs




- Divided SPC p -value:

Kolmogorov-Smirnov uniformity test p -value

$$p_{\text{dSPC}}(\mathbf{X}) := p_{\text{KS}}(p_{\text{SPC}}(\mathbf{x}^{(1)}), \dots, p_{\text{SPC}}(\mathbf{x}^{(k)}))$$

Divided SPCs



- Divided SPC p -value:

Kolmogorov-Smirnov uniformity test p -value

$$p_{\text{dSPC}}(\mathbf{X}) := p_{\text{KS}}(p_{\text{SPC}}(\mathbf{x}^{(1)}), \dots, p_{\text{SPC}}(\mathbf{x}^{(k)}))$$

Question 3: How is calibration and power of divided SPCs?

**Divided SPCs are well-calibrated and
have asymptotic power 1**

Divided SPCs are well-calibrated and have asymptotic power 1

Theorem [LH22]: if H_0 holds, then $\mathbb{P}\{p_{dSPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Divided SPCs are well-calibrated and have asymptotic power 1

Theorem [LH22]: if H_0 holds, then $\mathbb{P}\{p_{dSPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$ or $\rho \neq 1$ (so H_1 holds),
then power $\xrightarrow{P} 1$

Divided SPCs are well-calibrated and have asymptotic power 1

Theorem [LH22]: if H_0 holds, then $\mathbb{P}\{p_{dSPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Theorem [LH22]: if $\nu_{\circ} \neq \nu(\theta_{\star})$ or $\rho \neq 1$ (so H_1 holds),
then power $\xrightarrow{P} 1$

Question 4: How to scale k in divided SPCs?

Divided SPCs are well-calibrated and have asymptotic power 1

Theorem [LH22]: if H_0 holds, then $\mathbb{P}\{p_{dSPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Theorem [LH22]: if $\nu_\circ \neq \nu(\theta_\star)$ or $\rho \neq 1$ (so H_1 holds),
then power $\xrightarrow{P} 1$

Question 4: How to scale k in divided SPCs?

Condition: theorems hold only if $k = O(N^{\beta-\varepsilon})$ [for some $\varepsilon > 0$]

Divided SPCs are well-calibrated and have asymptotic power 1

Theorem [LH22]: if H_0 holds, then $\mathbb{P}\{p_{dSPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Theorem [LH22]: if $\nu_\circ \neq \nu(\theta_\star)$ or $\rho \neq 1$ (so H_1 holds),
then power $\xrightarrow{P} 1$

Question 4: How to scale k in divided SPCs? *$\beta = 1/2$ if model is smooth enough and higher-order moments exist*

Condition: theorems hold only if $k = O(N^{\beta-\varepsilon})$ [for some $\varepsilon > 0$]

Divided SPCs are well-calibrated and have asymptotic power 1

Theorem [LH22]: if H_0 holds, then $\mathbb{P}\{p_{dSPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Theorem [LH22]: if $\nu_\circ \neq \nu(\theta_\star)$ or $\rho \neq 1$ (so H_1 holds),
then power $\xrightarrow{P} 1$

Question 4: How to scale k in divided SPCs? *$\beta = 1/2$ if model is smooth enough and higher-order moments exist*

Condition: theorems hold only if $k = O(N^{\beta-\varepsilon})$ [for some $\varepsilon > 0$]

- In practice, we set $k = N^{0.49}$

Numerical investigation: subtle-to-mild misspecification

Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Pois}(\theta)$


Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
- Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$


Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
 - Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$
- over-dispersed*
- 


Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
 - Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$
- 
- The diagram shows two curved arrows. The first arrow, labeled *over-dispersed*, points from the text above to $\text{NegBin}(2, \tau)$ in the second bullet point. The second arrow, labeled *under-dispersed*, points from the text above to $\text{Binom}(30, p)$ in the same bullet point.


Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
 - Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$
 - For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, we have
- 
- The diagram shows two curved arrows. One arrow originates from the label *over-dispersed* and points to the $\text{NegBin}(2, \tau)$ distribution in the second bullet point. The other arrow originates from the label *under-dispersed* and points to the $\text{Binom}(30, p)$ distribution in the same bullet point.

Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
 - Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$
 - For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, we have
 - $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$
- 


Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
- Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$
 - over-dispersed*
 - under-dispersed*
- For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, we have
 - $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$
- Recall:


Numerical investigation: subtle-to-mild misspecification

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$


- Model $P_{\theta} = \text{Poiss}(\theta)$
- Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$ 
- For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, we have
 - $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$
- Recall:
 - Single SPCs: if $\nu_{\circ} = \nu(\theta_{\star})$ asymptotic power depends on ρ

Numerical investigation: subtle-to-mild misspecification

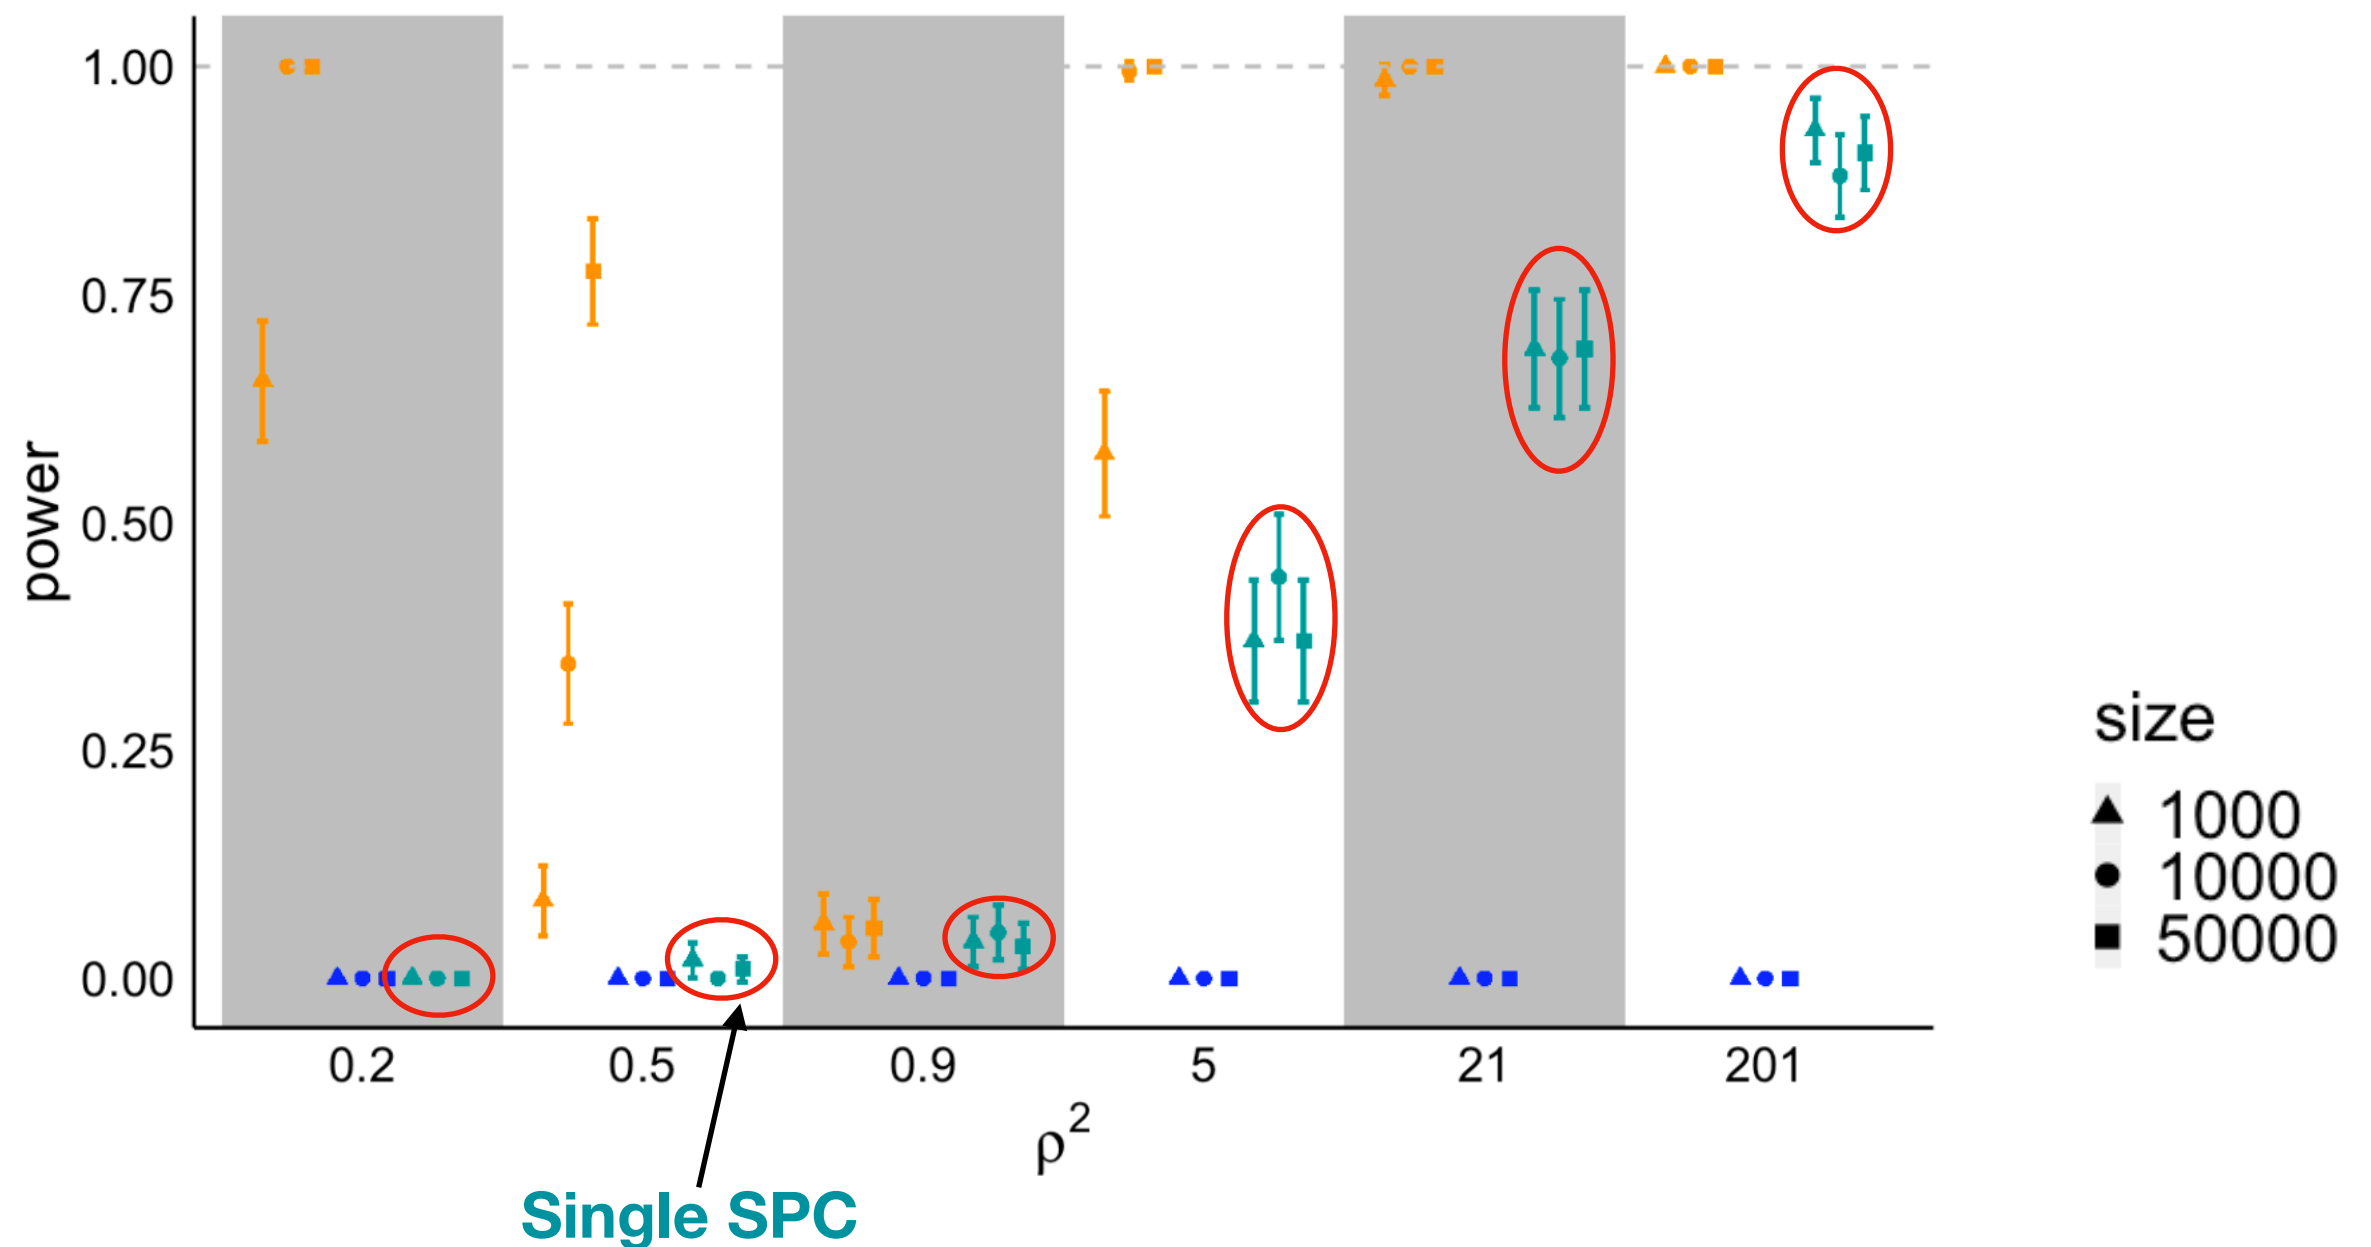
$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

- Model $P_{\theta} = \text{Poiss}(\theta)$
- Generate data from $P_{\circ} = \text{NegBin}(2, \tau)$ and $P_{\circ} = \text{Binom}(30, p)$ 
- For mean statistic $T(\mathbf{x}) = \sum_n x_n / N$, we have
 - $\nu_{\circ} = \nu(\theta_{\star})$ and $\rho = \sigma_{\circ} / \sigma(\theta_{\star})$
- Recall:
 - Single SPCs: if $\nu_{\circ} = \nu(\theta_{\star})$ asymptotic power depends on ρ
 - Divided SPCs: asymptotic power 1

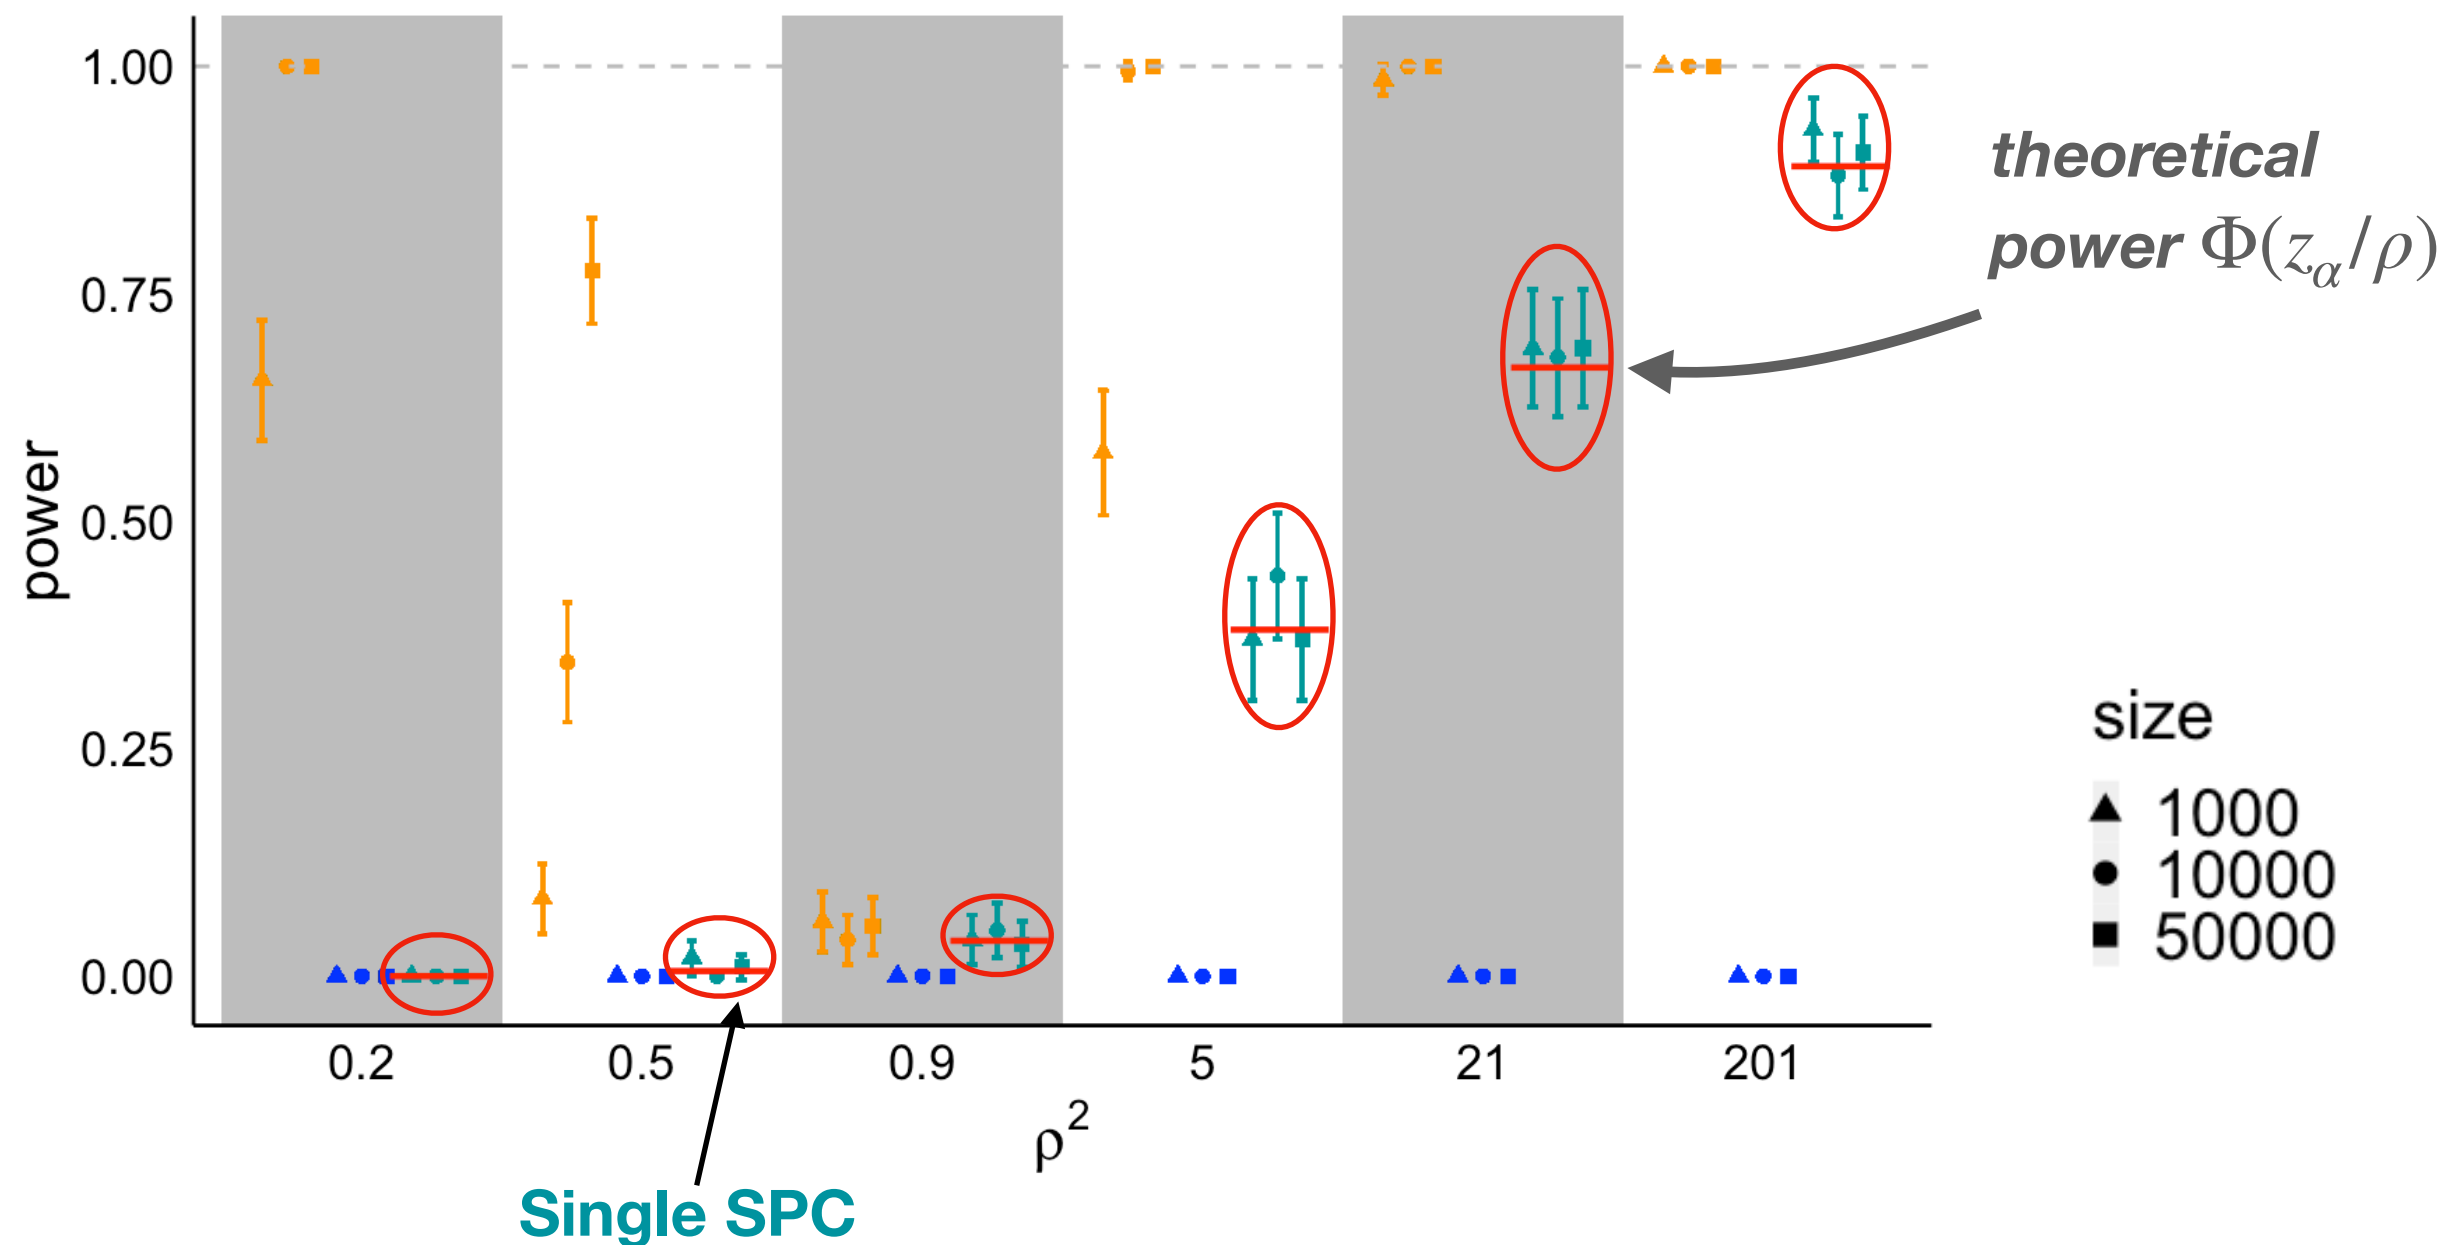
Divided SPCs have superior power



Divided SPCs have superior power




Divided SPCs have superior power



✓ If $\nu_\circ = \nu(\theta_\star)$ asymptotic power of single SPC depends on ρ

Divided SPCs have superior power



- ✓ If $\nu_o = \nu(\theta_\star)$ asymptotic power of single SPC depends on ρ
- ✓ Divided SPCs have asymptotic power 1

Comparison between single and divided SPCs

Comparison between single and divided SPCs

Single SPCs	Divided SPCs

Comparison between single and divided SPCs

Single SPCs	Divided SPCs
Compute posterior once	

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Asymptotic power could be small (if subtle misspecification)

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Asymptotic power could be small (if subtle misspecification)

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Asymptotic power = 1

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Asymptotic power could be small (if subtle misspecification)

Doesn't requires much data
(single “train-test” split)

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Asymptotic power = 1

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Asymptotic power could be small (if subtle misspecification)

Doesn't requires much data
(single "train-test" split)

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Asymptotic power = 1

Requires more data
(compute k single SPC p -values)

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Asymptotic power could be small (if subtle misspecification)

Doesn't requires much data
(single "train-test" split)

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Asymptotic power = 1

Requires more data
(compute k single SPC p -values)

-
- Single SPCs are always a good place to start

Comparison between single and divided SPCs

Single SPCs

Compute posterior once

Asymptotic power could be small (if subtle misspecification)

Doesn't requires much data
(single "train-test" split)

Divided SPCs

Compute posterior k times
(with k times smaller datasets)

Asymptotic power = 1

Requires more data
(compute k single SPC p -values)

-
- Single SPCs are always a good place to start
 - Consider divided SPCs if single SPCs inconclusive

How to “split” in single SPCs with time-series data?

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC

$$x_1 \quad \dots \quad x_{\lfloor N/2 \rfloor} \quad x_{\lfloor N/2 \rfloor + 1} \quad \dots \quad x_N$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC

$$x_1 \quad \dots \quad x_{\lfloor N/2 \rfloor} \quad \vdots \quad x_{\lfloor N/2 \rfloor + 1} \quad \dots \quad x_N$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC

$$\underbrace{x_1 \dots x_{[N/2]}}_{\mathbf{X}_{\text{obs}}} \vdots x_{[N/2]+1} \dots x_N$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC

$$\underbrace{x_1 \dots x_{\lfloor N/2 \rfloor}}_{\mathbf{X}_{\text{obs}}} : \underbrace{x_{\lfloor N/2 \rfloor + 1} \dots x_N}_{\mathbf{X}_{\text{ho}}}$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC



predict future well?

$$\underbrace{x_1 \dots x_{\lfloor N/2 \rfloor}}_{\mathbf{X}_{\text{obs}}} : \underbrace{x_{\lfloor N/2 \rfloor + 1} \dots x_N}_{\mathbf{X}_{\text{ho}}}$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC



predict future well?

interpolated single SPC

$$\underbrace{x_1 \dots x_{[N/2]}}_{\mathbf{X}_{\text{obs}}} : \underbrace{x_{[N/2]+1} \dots x_N}_{\mathbf{X}_{\text{ho}}}$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC



predict future well?

$$\underbrace{x_1 \quad \dots \quad x_{\lfloor N/2 \rfloor}}_{\mathbf{X}_{\text{obs}}} \vdots \underbrace{x_{\lfloor N/2 \rfloor + 1} \quad \dots \quad x_N}_{\mathbf{X}_{\text{ho}}}$$

interpolated single SPC

$$x_1 \quad x_2 \quad \dots \quad x_i \quad x_{i+1} \quad \dots \quad x_{N-1} \quad x_N$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC



predict future well?

$$\underbrace{x_1 \dots x_{\lfloor N/2 \rfloor}}_{\mathbf{X}_{\text{obs}}} : \underbrace{x_{\lfloor N/2 \rfloor + 1} \dots x_N}_{\mathbf{X}_{\text{ho}}}$$

interpolated single SPC

$$x_1 \quad x_2 \quad \dots \quad x_i \quad x_{i+1} \quad \dots \quad x_{N-1} \quad x_N$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC




predict future well?

$$\underbrace{x_1 \dots x_{\lfloor N/2 \rfloor}}_{\mathbf{X}_{\text{obs}}} : \underbrace{x_{\lfloor N/2 \rfloor + 1} \dots x_N}_{\mathbf{X}_{\text{ho}}}$$

interpolated single SPC

$$x_1 \quad x_2 \quad \dots \quad x_i \quad x_{i+1} \quad \dots \quad x_{N-1} \quad x_N$$



$$\underbrace{x_1 \quad x_3 \quad \dots \quad x_i \quad x_{i+2} \quad \dots \quad x_{N-1}}_{\mathbf{X}_{\text{obs}}} \quad \underbrace{x_2 \quad x_4 \quad \dots \quad x_{i+1} \quad x_{i+3} \quad \dots \quad x_N}_{\mathbf{X}_{\text{ho}}}$$

How to “split” in single SPCs with time-series data?

- If data are structured: must decide how to split the data
- Example: single 0.5-SPC with time-series data

extrapolated single SPC



predict future well?

$$\underbrace{x_1 \dots x_{\lfloor N/2 \rfloor}}_{\mathbf{X}_{\text{obs}}} \vdots \underbrace{x_{\lfloor N/2 \rfloor + 1} \dots x_N}_{\mathbf{X}_{\text{ho}}}$$

interpolated single SPC



explain data well?

$$x_1 \quad x_2 \quad \dots \quad x_i \quad x_{i+1} \quad \dots \quad x_{N-1} \quad x_N$$



$$\underbrace{x_1 \quad x_3 \quad \dots \quad x_i \quad x_{i+2} \quad \dots \quad x_{N-1}}_{\mathbf{X}_{\text{obs}}} \quad \underbrace{x_2 \quad x_4 \quad \dots \quad x_{i+1} \quad x_{i+3} \quad \dots \quad x_N}_{\mathbf{X}_{\text{ho}}}$$

SPCs for time series data

SPCs for time series data

- Airline delays data ($N = 327,346$)

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance
 - Response: arrival delays in minutes

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance
 - Response: arrival delays in minutes
 - Success rate statistic $T(\mathbf{x}) := N / \sum_n x_n$

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance
 - Response: arrival delays in minutes
 - Success rate statistic $T(\mathbf{x}) := N / \sum_n x_n$

*model fits
data well*



Methods	P-values
PPC	0.132
Interpolated single 0.5-SPC	0.626
Extrapolated single 0.5-SPC	0.000
Double-interpolated divided 0.5-SPC, $k = N^{0.49}$	0.000
Interpolated divided extrapolated 0.5-SPC, $k = N^{0.49}$	0.001

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance
 - Response: arrival delays in minutes
 - Success rate statistic $T(\mathbf{x}) := N / \sum_n x_n$

Methods	P-values
PPC	0.132
Interpolated single 0.5-SPC	0.626
Extrapolated single 0.5-SPC	0.000
Double-interpolated divided 0.5-SPC, $k = N^{0.49}$	0.000
Interpolated divided extrapolated 0.5-SPC, $k = N^{0.49}$	0.001

*model fits
data well*

*model does not
predict future well*

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance
 - Response: arrival delays in minutes
 - Success rate statistic $T(\mathbf{x}) := N / \sum_n x_n$

Methods	P-values
PPC	0.132
Interpolated single 0.5-SPC	0.626
Extrapolated single 0.5-SPC	0.000
Double-interpolated divided 0.5-SPC, $k = N^{0.49}$	0.000
Interpolated divided extrapolated 0.5-SPC, $k = N^{0.49}$	0.001

*model fits
data well*

*both **divided SPCs**
pick up misspecification*

*model **does not**
predict future well*

SPCs for time series data

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: month, day of week, distance
 - Response: arrival delays in minutes
 - Success rate statistic $T(\mathbf{x}) := N / \sum_n x_n$

Methods	P-values
PPC	0.132
Interpolated single 0.5-SPC	0.626
Extrapolated single 0.5-SPC	0.000
Double-interpolated divided 0.5-SPC, $k = N^{0.49}$	0.000
Interpolated divided extrapolated 0.5-SPC, $k = N^{0.49}$	0.001

*model fits
data well*

*both **divided SPCs**
pick up misspecification*

*model **does not**
predict future well*

- Different SPCs offer different (potentially useful!) insights

Concluding Thoughts

Concluding Thoughts

- ✓ Unlike PPCs, SPCs are **well-calibrated and well-powered...**

Concluding Thoughts

- ✓ Unlike PPCs, SPCs are **well-calibrated and well-powered...**
while also being general-purpose, computational efficient, and flexible

Concluding Thoughts

- ✓ Unlike PPCs, SPCs are **well-calibrated and well-powered...**
while also being general-purpose, computational efficient, and flexible
- ✓ Single SPCs and divided SPCs offer **complimentary strengths**

Concluding Thoughts

- ✓ Unlike PPCs, SPCs are **well-calibrated and well-powered...**
while also being general-purpose, computational efficient, and flexible
- ✓ Single SPCs and divided SPCs offer **complimentary strengths**
- ✓ SPCs introduce extra flexibility with different structured data and complex models such as time series, hierarchical, spatial data...

Concluding Thoughts

- ✓ Unlike PPCs, SPCs are **well-calibrated and well-powered...**
while also being general-purpose, computational efficient, and flexible
- ✓ Single SPCs and divided SPCs offer **complimentary strengths**
- ✓ SPCs introduce extra flexibility with different structured data and complex models such as time series, hierarchical, spatial data...
 - More to do: non-asymptotic results, non-normal limiting distributions, cross-validation versions, ...

Concluding Thoughts

- ✓ Unlike PPCs, SPCs are **well-calibrated and well-powered...** while also being general-purpose, computational efficient, and flexible
- ✓ Single SPCs and divided SPCs offer **complimentary strengths**
- ✓ SPCs introduce extra flexibility with different structured data and complex models such as time series, hierarchical, spatial data...
 - More to do: non-asymptotic results, non-normal limiting distributions, cross-validation versions, ...

J. Li and J. H. Huggins (2022). Calibrated Model Criticism Using Split Predictive Checks. *arXiv:2203.15897 [stat.ME]*

read the paper



Details of asymptotic theory for single SPCs

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

where $\Sigma_{\star} = \text{Cov}\{\nabla \ell(\theta_{\star}; X)\}$ and $J_{\star} = -\mathbb{E}\{\nabla^{\otimes 2} \ell(\theta_{\star}; X)\}$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

q controls the weights of two forms of mismatch

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

where $\Sigma_{\star} = \text{Cov}\{\nabla \ell(\theta_{\star}; X)\}$ and $J_{\star} = -\mathbb{E}\{\nabla^{\otimes 2} \ell(\theta_{\star}; X)\}$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

q controls the weights of two forms of mismatch

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

where $\Sigma_{\star} = \text{Cov}\{\nabla \ell(\theta_{\star}; X)\}$ and $J_{\star} = -\mathbb{E}\{\nabla^{\otimes 2} \ell(\theta_{\star}; X)\}$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

q controls the weights of two forms of mismatch

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

where $\Sigma_{\star} = \text{Cov}\{\nabla \ell(\theta_{\star}; X)\}$ and $J_{\star} = -\mathbb{E}\{\nabla^{\otimes 2} \ell(\theta_{\star}; X)\}$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$
- if model misspecified, then in general $\rho \neq 1$:

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

q controls the weights of two forms of mismatch

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

where $\Sigma_{\star} = \text{Cov}\{\nabla \ell(\theta_{\star}; X)\}$ and $J_{\star} = -\mathbb{E}\{\nabla^{\otimes 2} \ell(\theta_{\star}; X)\}$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$
- if model misspecified, then in general $\rho \neq 1$:
 - when uncertainty is under-estimated, $\rho \uparrow \infty$ and power $\uparrow 1$

Details of asymptotic theory for single SPCs

$$\nu_{\circ} := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\circ}(d\mathbf{x}), \quad \nu(\theta_{\star}) := \lim_{N \rightarrow \infty} \int T(\mathbf{x}) P_{\theta_{\star}}(d\mathbf{x})$$

Theorem: if $\nu_{\circ} = \nu(\theta_{\star})$, $\mathbb{P}\{p_{SPC}(\mathbf{X}) < \alpha\} \xrightarrow{P} \Phi\left(\frac{z_{\alpha}}{\rho}\right)$

q controls the weights of two forms of mismatch

$$\rho^2 := \frac{q\sigma_{\circ}^2 + (1-q)\dot{\nu}(\theta_{\star})^T \overbrace{J_{\star}^{-1} \Sigma_{\star} J_{\star}^{-1}}^{\text{🥪}} \dot{\nu}(\theta_{\star})}{q\sigma(\theta_{\star})^2 + (1-q)\dot{\nu}(\theta_{\star})^T J_{\star}^{-1} \dot{\nu}(\theta_{\star})}$$

where $\Sigma_{\star} = \text{Cov}\{\nabla \ell(\theta_{\star}; X)\}$ and $J_{\star} = -\mathbb{E}\{\nabla^{\otimes 2} \ell(\theta_{\star}; X)\}$

- if model correctly specified, then $\rho = 1$ and $\mathbb{P}\{p_{SPC}(\mathbf{x}) < \alpha\} \xrightarrow{P} \alpha$
- if model misspecified, then in general $\rho \neq 1$:
 - ▶ when uncertainty is under-estimated, $\rho \uparrow \infty$ and power $\uparrow 1$
 - ▶ when uncertainty is over-estimated, $\rho \downarrow 0$ and power $\downarrow 0$

Time series data using negative binomial regression model

Time series data using negative binomial regression model

- Airline delays data ($N = 327,346$)

Time series data using negative binomial regression model

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model

Time series data using negative binomial regression model


- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - Covariates: Month, Day of Week and Distance

Time series data using negative binomial regression model

- Airline delays data ($N = 327,346$)
- Consider negative binomial regression model
 - ▶ Covariates: Month, Day of Week and Distance
 - ▶ Response: Arrival delays in minutes


Time series data using negative binomial regression model

- Airline delays data (N = 327,346)
- Consider negative binomial regression model
 - Covariates: Month, Day of Week and Distance
 - Response: Arrival delays in minutes



Time series data using negative binomial regression model

- Airline delays data (N = 327,346)
- Consider negative binomial regression model
 - Covariates: Month, Day of Week and Distance
 - Response: Arrival delays in minutes



How to “split” in divided SPC with time-series data?

How to “split” in divided SPC with time-series data?

- Divided SPC

How to “split” in divided SPC with time-series data?

- Divided SPC

$$X_1 \quad \dots \quad \dots \quad X_i \quad \dots \quad \dots \quad X_N$$


How to “split” in divided SPC with time-series data?

- Divided SPC

$$X_1 \quad \cdot \cdot \cdot \quad \cdots \quad X_i \quad \cdot \cdot \cdot \quad \cdots \quad X_N$$


How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC




How to “split” in divided SPC with time-series data?

- Divided SPC





extrapolated dividing





How to “split” in divided SPC with time-series data?

- Divided SPC





How to “split” in divided SPC with time-series data?

- Divided SPC





How to “split” in divided SPC with time-series data?

- Divided SPC





How to “split” in divided SPC with time-series data?

- Divided SPC



How to “split” in divided SPC with time-series data?

- Divided SPC



How to “split” in divided SPC with time-series data?

- Divided SPC

