

# Likelihood-Free Inference with Generative Neural Networks via Scoring Rule Minimization

Lorenzo Pacchiardi

Department of Statistics, University of Oxford

*lorenzo.pacchiardi@stats.ox.ac.uk*

ISBA 2022, 30th June 2022

## Background: Likelihood-Free Inference

# Likelihood-Free Inference (LFI)

## Intractable-Likelihood model $P(\cdot|\theta)$

- can simulate data:

$$\mathbf{y} \sim P(\cdot|\theta), \mathbf{y} \in \mathcal{Y} \subseteq \mathbb{R}^d$$

- cannot evaluate  $p(\mathbf{y}|\theta)$

# Likelihood-Free Inference (LFI)

## Intractable-Likelihood model $P(\cdot|\theta)$

- can simulate data:

$$\mathbf{y} \sim P(\cdot|\theta), \mathbf{y} \in \mathcal{Y} \subseteq \mathbb{R}^d$$

- cannot evaluate  $p(\mathbf{y}|\theta)$

standard posterior  $\pi(\theta|\mathbf{y}) \propto \pi(\theta)p(\mathbf{y}|\theta)$  is inaccessible!

# Likelihood-Free Inference (LFI)

## Intractable-Likelihood model $P(\cdot|\theta)$

- can simulate data:

$$\mathbf{y} \sim P(\cdot|\theta), \mathbf{y} \in \mathcal{Y} \subseteq \mathbb{R}^d$$

- cannot evaluate  $p(\mathbf{y}|\theta)$

standard posterior  $\pi(\theta|\mathbf{y}) \propto \pi(\theta)p(\mathbf{y}|\theta)$  is inaccessible!

---


## Likelihood-Free Inference (LFI) approaches:

- Allow to sample from approximate posterior,
- rely on **drawing simulations** from the model.

# E.g.: Approximate Bayesian Computation

Iterate:


- 1 Draw  $\theta^{(j)} \sim \pi(\theta)$
- 2 Simulate a dataset  $\mathbf{y}_{\text{sim}} \sim P(\cdot | \theta^{(j)})$
- 3 Compute some **statistics**  $s(\mathbf{y})$  of the simulated and observed datasets
- 4 If **distance**  $d(s(\mathbf{y}_{\text{sim}}), s(\mathbf{y}_o)) \leq \epsilon$  (threshold)  $\implies$  accept  $\theta^{(j)}$ ; otherwise, reject



# E.g.: Approximate Bayesian Computation

Iterate:

- 1 Draw  $\theta^{(j)} \sim \pi(\theta)$
- 2 Simulate a dataset  $\mathbf{y}_{\text{sim}} \sim P(\cdot | \theta^{(j)})$
- 3 Compute some **statistics**  $s(\mathbf{y})$  of the simulated and observed datasets
- 4 If **distance**  $d(s(\mathbf{y}_{\text{sim}}), s(\mathbf{y}_o)) \leq \epsilon$  (threshold)  $\implies$  accept  $\theta^{(j)}$ ; otherwise, reject



Accepted  $\theta^{(j)}$ 's are distributed according to the ABC posterior:

$$\pi^\epsilon(\theta | s(\mathbf{y}_o)) \propto \pi(\theta) \int \mathbb{1}[d(s(\mathbf{y}_o), s(\mathbf{y}_{\text{sim}})) \leq \epsilon] p(\mathbf{y}_{\text{sim}} | \theta) d\mathbf{y}_{\text{sim}}$$

## Background: conditional generative networks



# Conditional generative networks

Defined by:

- a neural network  $f_\phi : \mathcal{Z} \times \mathcal{Y} \rightarrow \Theta$ ;
- a probability distribution  $P_{\mathbf{z}}$  over the space  $\mathcal{Z}$ .

# Conditional generative networks


Defined by:

- a neural network  $f_\phi : \mathcal{Z} \times \mathcal{Y} \rightarrow \Theta$ ;
- a probability distribution  $P_z$  over the space  $\mathcal{Z}$ .

These induce a distribution  $Q_\phi(\cdot|\mathbf{y})$  over  $\Theta$  by:

- 1 Sampling  $\mathbf{z} \sim P_z$ ,
- 2 computing  $f_\phi(\mathbf{z}, \mathbf{y})$ .

In statistical notation:  $Q_\phi(\cdot|\mathbf{y}) = \underbrace{f_\phi(\cdot, \mathbf{y})\#P_z}_{\text{"pushforward"}}$ .



# Conditional generative networks


Defined by:

- a neural network  $f_\phi : \mathcal{Z} \times \mathcal{Y} \rightarrow \Theta$ ;
- a probability distribution  $P_z$  over the space  $\mathcal{Z}$ .

These induce a distribution  $Q_\phi(\cdot|\mathbf{y})$  over  $\Theta$  by:

- 1 Sampling  $\mathbf{z} \sim P_z$ ,
- 2 computing  $f_\phi(\mathbf{z}, \mathbf{y})$ .

In statistical notation:  $Q_\phi(\cdot|\mathbf{y}) = \underbrace{f_\phi(\cdot, \mathbf{y})\#P_z}_{\text{"pushforward"}}$ .



**Idea:**

- Use a **generative network to approximate the posterior**
- Can **easily obtain samples** from it (no MCMC)

# Conditional generative networks


Defined by:

- a neural network  $f_\phi : \mathcal{Z} \times \mathcal{Y} \rightarrow \Theta$ ;
- a probability distribution  $P_z$  over the space  $\mathcal{Z}$ .

These induce a distribution  $Q_\phi(\cdot|\mathbf{y})$  over  $\Theta$  by:

- 1 Sampling  $\mathbf{z} \sim P_z$ ,
- 2 computing  $f_\phi(\mathbf{z}, \mathbf{y})$ .

In statistical notation:  $Q_\phi(\cdot|\mathbf{y}) = \underbrace{f_\phi(\cdot, \mathbf{y})\#P_z}_{\text{"pushforward"}}$



**Idea:**

- Use a **generative network to approximate the posterior**
- Can **easily obtain samples** from it (no MCMC)
- But you **can't evaluate the density**: how can you train it?

# Posterior inference via Generative Adversarial Networks (Ramesh et al., 2022)

# Posterior via Generative Adversarial Networks (GANs)

(Ramesh et al., 2022)

- Consider a **discriminator** neural network  $D_\psi : \Theta \times \mathcal{Y} \rightarrow [0, 1]$
- Define loss (Goodfellow et al., 2014):

$$L(\phi, \psi) := \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_\psi(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_\phi(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_\psi(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

# Posterior via Generative Adversarial Networks (GANs)

(Ramesh et al., 2022)

- Consider a **discriminator** neural network  $D_\psi : \Theta \times \mathcal{Y} \rightarrow [0, 1]$
- Define loss (Goodfellow et al., 2014):

$$L(\phi, \psi) := \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_\psi(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_\phi(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_\psi(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

- Solving  $\min_\phi \max_\psi L(\phi, \psi) \implies Q_\phi(\cdot | \mathbf{y}) = \Pi(\cdot | \mathbf{y}) \forall \mathbf{y} : p(\mathbf{y}) > 0$

# Posterior via Generative Adversarial Networks (GANs)

(Ramesh et al., 2022)

- Consider a **discriminator** neural network  $D_\psi : \Theta \times \mathcal{Y} \rightarrow [0, 1]$
- Define loss (Goodfellow et al., 2014):

$$L(\phi, \psi) := \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_\psi(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_\phi(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_\psi(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

- Solving  $\min_\phi \max_\psi L(\phi, \psi) \implies Q_\phi(\cdot | \mathbf{y}) = \Pi(\cdot | \mathbf{y}) \forall \mathbf{y} : p(\mathbf{y}) > 0$



# Adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

- 1 Generate a dataset of **parameter-simulations pairs** from the likelihood-free model:

$$(\boldsymbol{\theta}_i, \mathbf{y}_i)_{i=1}^n, \boldsymbol{\theta}_i \sim \Pi \text{ and } \mathbf{y}_i \sim P(\cdot | \boldsymbol{\theta}_i)$$

# Adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

- 1 Generate a dataset of **parameter-simulations pairs** from the likelihood-free model:

$$(\boldsymbol{\theta}_i, \mathbf{y}_i)_{i=1}^n, \boldsymbol{\theta}_i \sim \Pi \text{ and } \mathbf{y}_i \sim P(\cdot | \boldsymbol{\theta}_i) \iff \mathbf{y}_i \sim P \text{ and } \boldsymbol{\theta}_i \sim \Pi(\cdot | \mathbf{y}_i)$$

# Adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

- ① Generate a dataset of **parameter-simulations pairs** from the likelihood-free model:

$$(\boldsymbol{\theta}_i, \mathbf{y}_i)_{i=1}^n, \boldsymbol{\theta}_i \sim \Pi \text{ and } \mathbf{y}_i \sim P(\cdot | \boldsymbol{\theta}_i) \iff \mathbf{y}_i \sim P \text{ and } \boldsymbol{\theta}_i \sim \Pi(\cdot | \mathbf{y}_i)$$

- ② Train by **alternating** stochastic gradient ascent/descent over  $\psi$  and  $\phi$ :

**Require:** Generative net  $f_{\phi}$ , discriminator  $D_{\psi}$ , learning rates  $\epsilon, \gamma$ .

- 1: **for** each training pair  $(\boldsymbol{\theta}_i, \mathbf{y}_i)$  **do**
- 2:     Sample  $\mathbf{z}_i \sim P_{\mathbf{z}}$
- 3:     Obtain  $\tilde{\boldsymbol{\theta}}_i^{\phi} = f_{\phi}(\mathbf{z}_i, \mathbf{y}_i)$
- 4:     Set  $\psi \leftarrow \psi + \gamma \cdot \nabla_{\psi} \left[ \log D_{\psi}(\boldsymbol{\theta}_i, \mathbf{y}_i) + \log(1 - D_{\psi}(\tilde{\boldsymbol{\theta}}_i^{\phi}, \mathbf{y}_i)) \right]$
- 5:     Set  $\phi \leftarrow \phi - \epsilon \cdot \nabla_{\phi} \left[ \log(1 - D_{\psi}(\tilde{\boldsymbol{\theta}}_i^{\phi}, \mathbf{y}_i)) \right]$
- 6: **end for**

- Use batch of training data to estimate  $\mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} \dots$
- and **draws from the generative network** to estimate  $\mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})}$

# Adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_{\psi}(\tilde{\boldsymbol{\theta}}, \mathbf{Y}) \right) \right) \right]$$

- ① Generate a dataset of **parameter-simulations pairs** from the likelihood-free model:

$$(\boldsymbol{\theta}_i, \mathbf{y}_i)_{i=1}^n, \boldsymbol{\theta}_i \sim \Pi \text{ and } \mathbf{y}_i \sim P(\cdot | \boldsymbol{\theta}_i) \iff \mathbf{y}_i \sim P \text{ and } \boldsymbol{\theta}_i \sim \Pi(\cdot | \mathbf{y}_i)$$

- ② Train by **alternating** stochastic gradient ascent/descent over  $\psi$  and  $\phi$ :

**Require:** Generative net  $f_{\phi}$ , discriminator  $D_{\psi}$ , learning rates  $\epsilon, \gamma$ .

1: **for** each training pair  $(\boldsymbol{\theta}_i, \mathbf{y}_i)$  **do**

2:   Sample  $\mathbf{z}_i \sim P_{\mathbf{z}}$

3:   Obtain  $\tilde{\boldsymbol{\theta}}_i^{\phi} = f_{\phi}(\mathbf{z}_i, \mathbf{y}_i)$

4:   Set  $\psi \leftarrow \psi + \gamma \cdot \nabla_{\psi} \left[ \log D_{\psi}(\boldsymbol{\theta}_i, \mathbf{y}_i) + \log(1 - D_{\psi}(\tilde{\boldsymbol{\theta}}_i^{\phi}, \mathbf{y}_i)) \right]$

5:   Set  $\phi \leftarrow \phi - \epsilon \cdot \nabla_{\phi} \left[ \log(1 - D_{\psi}(\tilde{\boldsymbol{\theta}}_i^{\phi}, \mathbf{y}_i)) \right]$

6: **end for**

- Use batch of training data to estimate  $\mathbb{E}_{\boldsymbol{\theta} \sim \Pi(\cdot | \mathbf{Y})} \dots$
- and **draws from the generative network** to estimate  $\mathbb{E}_{\tilde{\boldsymbol{\theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})}$

- ③ For a real observation  $\mathbf{y}_o$ , you can directly get samples from  $\Pi(\cdot | \mathbf{y}_o)$ .

# Adversarial training

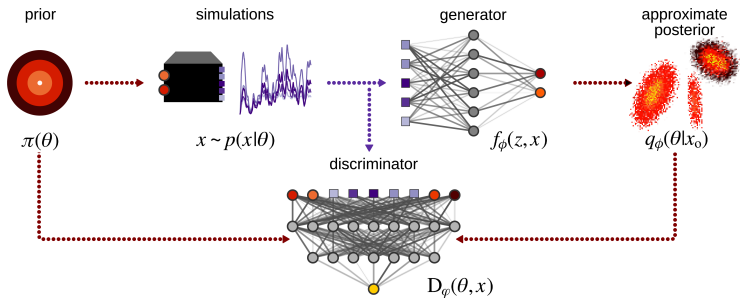


Figure: From Ramesh et al. (2022)

# Issues with adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} (\log (1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y})) \right]$$

# Issues with adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} (\log (1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y})) \right]$$

Issues:

- 1 **Unstable** min-max problem.

# Issues with adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} \left( \log \left( 1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y}) \right) \right) \right]$$

Issues:

- ① **Unstable** min-max problem.
- ② **biased gradients** w.r.t.  $\phi$  when:
  - doing few updates for  $\psi$
  - and using it to estimate gradients w.r.t.  $\phi$



# Issues with adversarial training

$$\min_{\phi} \max_{\psi} \mathbb{E}_{\mathbf{Y} \sim P} \left[ \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} (\log D_{\psi}(\boldsymbol{\Theta}, \mathbf{Y})) + \mathbb{E}_{\tilde{\boldsymbol{\Theta}} \sim Q_{\phi}(\cdot | \mathbf{Y})} \left( \log (1 - D_{\psi}(\tilde{\boldsymbol{\Theta}}, \mathbf{Y})) \right) \right]$$

Issues:

- ① **Unstable** min-max problem.
- ② **biased gradients** w.r.t.  $\phi$  when:
  - doing few updates for  $\psi$
  - and using it to estimate gradients w.r.t.  $\phi$
- ③ **Mode collapse.**

# Posterior inference via Scoring Rules Minimization for Generative Networks

$S(Q_\phi, \theta) \rightarrow$  “*penalty* when stating a distribution  $Q_\phi$  for a realization  $\theta$ ”.

$S(Q_\phi, \theta) \rightarrow$  “*penalty* when stating a distribution  $Q_\phi$  for a realization  $\theta$ ”.

---

**Expected Scoring Rule:** if  $\Theta \sim \Pi$ :

$$S(Q_\phi, \Pi) := \mathbb{E}_{\Theta \sim \Pi} S(Q_\phi, \Theta).$$

$S(Q_\phi, \theta) \rightarrow$  “*penalty* when stating a distribution  $Q_\phi$  for a realization  $\theta$ ”.

---

**Expected Scoring Rule:** if  $\Theta \sim \Pi$ :

$$S(Q_\phi, \Pi) := \mathbb{E}_{\Theta \sim \Pi} S(Q_\phi, \Theta).$$

**Strictly proper**  $S$ :  $S(Q_\phi, \Pi)$  is uniquely minimized by  $Q_\phi = \Pi$ .

# Some strictly proper scoring rules

## Energy Score:

$$S_E^{(\beta)}(Q_\phi, \theta) = 2 \cdot \mathbb{E} \left[ \|\tilde{\theta} - \theta\|_2^\beta \right] - \mathbb{E} \left[ \|\tilde{\theta} - \tilde{\theta}'\|_2^\beta \right], \quad \tilde{\theta}, \tilde{\theta}' \sim Q_\phi,$$

## Kernel Score: related to MMD<sup>2</sup>

$$S_k(Q_\phi, \theta) = \mathbb{E}[k(\tilde{\theta}, \tilde{\theta}')] - 2 \cdot \mathbb{E}[k(\tilde{\theta}, \theta)], \quad \tilde{\theta}, \tilde{\theta}' \sim Q_\phi.$$

$k \rightarrow$  positive definite kernel.

# Conditional Generative Networks via Scoring Rule Minimization

- Define:

$$J(\phi) := \mathbb{E}_{\mathbf{Y} \sim P} \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot | \mathbf{Y})} S(Q_\phi(\cdot | \mathbf{Y}), \boldsymbol{\Theta})$$

- For **strictly proper**  $S$ ,  $\min_\phi J(\phi) \implies Q_\phi(\cdot | \mathbf{y}) = \Pi(\cdot | \mathbf{y}) \ \forall \ \mathbf{y} : p(\mathbf{y}) > 0$ .

# Conditional Generative Networks via Scoring Rule Minimization

- Define:

$$J(\phi) := \mathbb{E}_{\mathbf{Y} \sim P} \mathbb{E}_{\boldsymbol{\theta} \sim \Pi(\cdot|\mathbf{Y})} S(Q_\phi(\cdot|\mathbf{Y}), \boldsymbol{\theta})$$

- For **strictly proper**  $S$ ,  $\min_{\phi} J(\phi) \implies Q_\phi(\cdot|\mathbf{y}) = \Pi(\cdot|\mathbf{y}) \forall \mathbf{y} : p(\mathbf{y}) > 0$ .

**Unbiased empirical estimate:**

$$\arg \min_{\phi} \left[ \hat{J}(\phi) := \frac{1}{n} \sum_{i=1}^n S(Q_\phi(\cdot|\mathbf{y}_i), \boldsymbol{\theta}_i) \right], \quad \boldsymbol{\theta}_i \sim \Pi \text{ and } \mathbf{y}_i \sim P(\cdot|\boldsymbol{\theta}_i),$$



# Conditional Generative Networks via Scoring Rule Minimization

- Define:

$$J(\phi) := \mathbb{E}_{\mathbf{Y} \sim P} \mathbb{E}_{\boldsymbol{\Theta} \sim \Pi(\cdot|\mathbf{Y})} S(Q_\phi(\cdot|\mathbf{Y}), \boldsymbol{\Theta})$$

- For **strictly proper**  $S$ ,  $\min_{\phi} J(\phi) \implies Q_\phi(\cdot|\mathbf{y}) = \Pi(\cdot|\mathbf{y}) \forall \mathbf{y} : p(\mathbf{y}) > 0$ .

**Unbiased empirical estimate:**

$$\arg \min_{\phi} \left[ \hat{J}(\phi) := \frac{1}{n} \sum_{i=1}^n S(Q_\phi(\cdot|\mathbf{y}_i), \boldsymbol{\theta}_i) \right], \quad \boldsymbol{\theta}_i \sim \Pi \text{ and } \mathbf{y}_i \sim P(\cdot|\boldsymbol{\theta}_i),$$

$\implies$  Unbiased estimate of  $\nabla_{\phi} S(Q_\phi(\cdot|\mathbf{y}_i), \boldsymbol{\theta}_i)$  **is enough** to train via SGD.

# How to get unbiased gradient estimates

We **need unbiased estimate** of  $\nabla_{\phi} S(Q_{\phi}(\cdot|\mathbf{y}), \theta)$ .

# How to get unbiased gradient estimates

We **need unbiased estimate** of  $\nabla_{\phi} S(Q_{\phi}(\cdot|\mathbf{y}), \theta)$ .

- We take  $S(Q_{\phi}(\cdot|\mathbf{y}), \theta) = \mathbb{E}_{\tilde{\theta}, \tilde{\theta}' \sim Q_{\phi}(\cdot|\mathbf{y})} \left[ g(\tilde{\theta}, \tilde{\theta}', \theta) \right]$  for some function  $g$

# How to get unbiased gradient estimates

We **need unbiased estimate** of  $\nabla_{\phi} S(Q_{\phi}(\cdot|\mathbf{y}), \theta)$ .

- We take  $S(Q_{\phi}(\cdot|\mathbf{y}), \theta) = \mathbb{E}_{\tilde{\theta}, \tilde{\theta}' \sim Q_{\phi}(\cdot|\mathbf{y})} [g(\tilde{\theta}, \tilde{\theta}', \theta)]$  for some function  $g \implies$  we can get **unbiased estimator from draws**  $\tilde{\theta}_j^{\phi} = f_{\phi}(\mathbf{Z}_j, \mathbf{y}) \sim Q_{\phi}(\cdot|\mathbf{y})$ ,  $j = 1, \dots, m$ :

$$\hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \frac{1}{m(m-1)} \sum_{j \neq k} g(\tilde{\theta}_j^{\phi}, \tilde{\theta}_k^{\phi}, \theta)$$

# How to get unbiased gradient estimates

We **need unbiased estimate** of  $\nabla_{\phi} S(Q_{\phi}(\cdot|\mathbf{y}), \theta)$ .

- We take  $S(Q_{\phi}(\cdot|\mathbf{y}), \theta) = \mathbb{E}_{\tilde{\theta}, \tilde{\theta}' \sim Q_{\phi}(\cdot|\mathbf{y})} [g(\tilde{\theta}, \tilde{\theta}', \theta)]$  for some function  $g \implies$  we can get **unbiased estimator from draws**  $\tilde{\theta}_j^{\phi} = f_{\phi}(\mathbf{Z}_j, \mathbf{y}) \sim Q_{\phi}(\cdot|\mathbf{y})$ ,  $j = 1, \dots, m$ :

$$\hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \frac{1}{m(m-1)} \sum_{j \neq k} g(\tilde{\theta}_j^{\phi}, \tilde{\theta}_k^{\phi}, \theta)$$

- With autodifferentiation: obtain  $\nabla_{\phi} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} \hat{S}(\{f_{\phi}(\mathbf{Z}_j, \mathbf{y})\}_j, \theta)$ .

# How to get unbiased gradient estimates

We **need unbiased estimate** of  $\nabla_{\phi} S(Q_{\phi}(\cdot|\mathbf{y}), \theta)$ .

- We take  $S(Q_{\phi}(\cdot|\mathbf{y}), \theta) = \mathbb{E}_{\tilde{\theta}, \tilde{\theta}' \sim Q_{\phi}(\cdot|\mathbf{y})} [g(\tilde{\theta}, \tilde{\theta}', \theta)]$  for some function  $g \implies$  we can get **unbiased estimator from draws**  $\tilde{\theta}_j^{\phi} = f_{\phi}(\mathbf{Z}_j, \mathbf{y}) \sim Q_{\phi}(\cdot|\mathbf{y})$ ,  $j = 1, \dots, m$ :

$$\hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \frac{1}{m(m-1)} \sum_{j \neq k} g(\tilde{\theta}_j^{\phi}, \tilde{\theta}_k^{\phi}, \theta)$$

- With autodifferentiation: obtain  $\nabla_{\phi} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} \hat{S}(\{f_{\phi}(\mathbf{Z}_j, \mathbf{y})\}_j, \theta)$ .
- But (swapping expectation and gradient):

$$\mathbb{E} \nabla_{\phi} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} \mathbb{E} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} S(Q_{\phi}, \theta)$$

# How to get unbiased gradient estimates

We **need unbiased estimate** of  $\nabla_{\phi} S(Q_{\phi}(\cdot|\mathbf{y}), \theta)$ .

- We take  $S(Q_{\phi}(\cdot|\mathbf{y}), \theta) = \mathbb{E}_{\tilde{\theta}, \tilde{\theta}' \sim Q_{\phi}(\cdot|\mathbf{y})} [g(\tilde{\theta}, \tilde{\theta}', \theta)]$  for some function  $g \implies$  we can get **unbiased estimator from draws**  $\tilde{\theta}_j^{\phi} = f_{\phi}(\mathbf{Z}_j, \mathbf{y}) \sim Q_{\phi}(\cdot|\mathbf{y}), j = 1, \dots, m$ :

$$\hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \frac{1}{m(m-1)} \sum_{j \neq k} g(\tilde{\theta}_j^{\phi}, \tilde{\theta}_k^{\phi}, \theta)$$

- With autodifferentiation: obtain  $\nabla_{\phi} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} \hat{S}(\{f_{\phi}(\mathbf{Z}_j, \mathbf{y})\}_j, \theta)$ .
- But (swapping expectation and gradient):

$$\mathbb{E} \nabla_{\phi} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} \mathbb{E} \hat{S}(\{\tilde{\theta}_j^{\phi}\}_j, \theta) = \nabla_{\phi} S(Q_{\phi}, \theta)$$

$\nabla_{\phi}$  of unbiased estimate of  $S \equiv$  unbiased estimate of  $\nabla_{\phi} S$

## SR minimization

**Require:** Generative net  $f_\phi$ , SR  $S$ , learning rate  $\epsilon$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:
- 3:
- 4:
- 5:
- 6: **end for**

---

## Adversarial training

**Require:** Generative net  $f_\phi$ , discriminator  $D_\psi$ , learning rates  $\epsilon, \gamma$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:
- 3:
- 4:
- 5:
- 6: **end for**



## SR minimization

**Require:** Generative net  $f_\phi$ , SR  $S$ , learning rate  $\epsilon$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample **multiple**  $\mathbf{z}_{i,1}, \dots, \mathbf{z}_{i,m} \sim P_{\mathbf{z}}$
- 3:
- 4:
- 5:
- 6: **end for**

---

## Adversarial training

**Require:** Generative net  $f_\phi$ , discriminator  $D_\psi$ , learning rates  $\epsilon, \gamma$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample  $\mathbf{z}_i \sim P_{\mathbf{z}}$
- 3:
- 4:
- 5:
- 6: **end for**

## SR minimization

**Require:** Generative net  $f_\phi$ , SR  $S$ , learning rate  $\epsilon$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample **multiple**  $\mathbf{z}_{i,1}, \dots, \mathbf{z}_{i,m} \sim P_{\mathbf{z}}$
- 3:     Obtain  $\tilde{\theta}_{i,j}^\phi = f_\phi(\mathbf{z}_{i,j}, \mathbf{y}_i)$
- 4:
- 5:
- 6: **end for**

---

## Adversarial training

**Require:** Generative net  $f_\phi$ , discriminator  $D_\psi$ , learning rates  $\epsilon, \gamma$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample  $\mathbf{z}_i \sim P_{\mathbf{z}}$
- 3:     Obtain  $\tilde{\theta}_i^\phi = f_\phi(\mathbf{z}_i, \mathbf{y}_i)$
- 4:
- 5:
- 6: **end for**

## SR minimization

**Require:** Generative net  $f_\phi$ , SR  $S$ , learning rate  $\epsilon$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
  - 2:     Sample **multiple**  $\mathbf{z}_{i,1}, \dots, \mathbf{z}_{i,m} \sim P_{\mathbf{z}}$
  - 3:     Obtain  $\tilde{\theta}_{i,j}^\phi = f_\phi(\mathbf{z}_{i,j}, \mathbf{y}_i)$
  - 4:     Compute unbiased estimate  $\hat{S}(\{\theta_{i,j}^\phi\}_j, \theta_i)$  from  $\tilde{\theta}_{i,j}^\phi$
  - 5:     Set  $\phi \leftarrow \phi - \epsilon \cdot \nabla_\phi \hat{S}(\{\theta_{i,j}^\phi\}_j, \theta_i)$
  - 6: **end for**
- 

## Adversarial training

**Require:** Generative net  $f_\phi$ , discriminator  $D_\psi$ , learning rates  $\epsilon, \gamma$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample  $\mathbf{z}_i \sim P_{\mathbf{z}}$
- 3:     Obtain  $\tilde{\theta}_i^\phi = f_\phi(\mathbf{z}_i, \mathbf{y}_i)$
- 4:     Set  $\psi \leftarrow \psi + \gamma \cdot \nabla_\psi \left[ \log D_\psi(\theta_i, \mathbf{y}_i) + \log(1 - D_\psi(\tilde{\theta}_i^\phi, \mathbf{y}_i)) \right]$
- 5:     Set  $\phi \leftarrow \phi - \epsilon \cdot \nabla_\phi \left[ \log(1 - D_\psi(\tilde{\theta}_i^\phi, \mathbf{y}_i)) \right]$
- 6: **end for**

## SR minimization

**Require:** Generative net  $f_\phi$ , SR  $S$ , learning rate  $\epsilon$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample **multiple**  $\mathbf{z}_{i,1}, \dots, \mathbf{z}_{i,m} \sim P_{\mathbf{z}}$
- 3:     Obtain  $\tilde{\theta}_{i,j}^\phi = f_\phi(\mathbf{z}_{i,j}, \mathbf{y}_i)$
- 4:     Compute unbiased estimate  $\hat{S}(\{\theta_{i,j}^\phi\}_j, \theta_i)$  from  $\tilde{\theta}_{i,j}^\phi$
- 5:     Set  $\phi \leftarrow \phi - \epsilon \cdot \nabla_\phi \hat{S}(\{\theta_{i,j}^\phi\}_j, \theta_i)$
- 6: **end for**

---

## Adversarial training

**Require:** Generative net  $f_\phi$ , discriminator  $D_\psi$ , learning rates  $\epsilon, \gamma$ .

- 1: **for** each training pair  $(\theta_i, \mathbf{y}_i)$  **do**
- 2:     Sample  $\mathbf{z}_i \sim P_{\mathbf{z}}$
- 3:     Obtain  $\tilde{\theta}_i^\phi = f_\phi(\mathbf{z}_i, \mathbf{y}_i)$
- 4:     Set  $\psi \leftarrow \psi + \gamma \cdot \nabla_\psi \left[ \log D_\psi(\theta_i, \mathbf{y}_i) + \log(1 - D_\psi(\tilde{\theta}_i^\phi, \mathbf{y}_i)) \right]$
- 5:     Set  $\phi \leftarrow \phi - \epsilon \cdot \nabla_\phi \left[ \log(1 - D_\psi(\tilde{\theta}_i^\phi, \mathbf{y}_i)) \right]$
- 6: **end for**


Simpler training than GAN:

- **no discriminator**  $D_\psi$
- **no min-max** objective

Need multiple simulations from  $Q_\phi$ , but **good results with as little as 3**.

# Simulations: noisy camera model

$\theta \in \mathbb{R}^{28 \times 28} \rightarrow$  EMNIST dataset,  $y_o \in \mathbb{R}^{28 \times 28}$  generated from  $\theta$  by blurring.



# Simulations: noisy camera model

$\theta \in \mathbb{R}^{28 \times 28} \rightarrow$  EMNIST dataset,  $y_o \in \mathbb{R}^{28 \times 28}$  generated from  $\theta$  by blurring.




Table: Training time (secs)

GAN	45398
Energy	22633
Kernel	22545

- **Bayesian Likelihood-Free Inference with generative networks used to approximate the posterior.**
- Alternative training method which...
  - ... is **faster** than adversarial,
  - does **not suffer from instabilities** of min-max training
  - and leads to **better results**.

L. Pacchiardi and R. Dutta. Likelihood-free inference with generative neural networks via scoring rule minimization. *arXiv preprint arXiv:2205.15784*, 2022.



---

## References

- I. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio. Generative adversarial nets. *Advances in neural information processing systems*, 27, 2014.
- L. Pacchiardi, R. Adewoyin, P. Dueben, and R. Dutta. Probabilistic forecasting with conditional generative networks via scoring rule minimization. *arXiv preprint arXiv:2112.08217*, 2022.
- P. Ramesh, J.-M. Lueckmann, J. Boelts, Á. Tejero-Cantero, D. S. Greenberg, P. J. Goncalves, and J. H. Macke. GATSBI: Generative adversarial training for simulation-based inference. In *International Conference on Learning Representations*, 2022.

---

## Contacts

website: [www.lorenzopacchiardi.me](http://www.lorenzopacchiardi.me)

email: [lorenzo.pacchiardi@stats.ox.ac.uk](mailto:lorenzo.pacchiardi@stats.ox.ac.uk)



## Backup slides

- If  $\mathbf{X}$  has some structure (say, it is on a 1D or 2D grid) raw SR discard that information.
- $\implies$  compute the SR on localized *patches* across the grid and cumulate the score; in this way, short-scale correlations are given more importance.
- The resulting SR is non-strictly proper  $\implies$  add the global SR to make it strictly proper.

The patched SR is:

$$S_p(P, \mathbf{x}) = w_1 S(P, \mathbf{x}) + w_2 \sum_{p \in \mathcal{P}} S(P|_p, \mathbf{x}|_p),$$

where  $w_1, w_2 > 0$ ,  $|_p$  denotes the restriction of a distribution or of a vector to a patch  $p$  and  $\mathcal{P}$  is a set of patches.

# Connection with normalizing flows

- Normalizing flows = generative networks with invertible  $f_\phi(\mathbf{z}, \mathbf{y})$  with respect to  $\mathbf{z}$ .
- Density evaluation is possible via change-of-variables formula  $\implies \phi$  is usually trained via maximum likelihood; e.g.:

$$\begin{aligned} & \underset{\phi}{\operatorname{argmin}} \mathbb{E}_{\mathbf{Y} \sim P} [\mathbb{KL}(\Pi(\cdot | \mathbf{Y}) \| Q_\phi(\cdot | \mathbf{Y}))] \\ &= \underset{\phi}{\operatorname{argmin}} \mathbb{E}_{\mathbf{Y} \sim P} \mathbb{E}_{\boldsymbol{\theta} \sim \Pi(\cdot | \mathbf{Y})} [-\log q_\phi(\boldsymbol{\theta} | \mathbf{Y})] \\ &= \underset{\phi}{\operatorname{argmin}} \mathbb{E}_{\boldsymbol{\theta} \sim \Pi} \mathbb{E}_{\mathbf{Y} \sim P(\cdot | \boldsymbol{\theta})} [-\log q_\phi(\boldsymbol{\theta} | \mathbf{Y})], \end{aligned}$$

which corresponds to our SR-based approach by identifying  $S(Q_\phi(\cdot | \mathbf{y}), \boldsymbol{\theta}) = -\log q_\phi(\boldsymbol{\theta} | \mathbf{y})$ , which is the strictly-proper logarithmic scoring rule.

## Additional results for noisy camera model

**Table:** Noisy Camera model: performance metrics, runtime and early stopping epoch for GAN and for the Energy and Kernel Score with patch size 8 and step 5. The latter methods achieved better performance with shorter training time. All methods are trained on a single GPU.

	RMSE ↓	Cal. Err. ↓	$R^2$ ↑	Runtime (sec)	Early stopping epoch
GAN	$0.25 \pm 0.19$	$0.50 \pm 0.00$	$-23.94 \pm 366.08$	45398	3600
Energy	$0.06 \pm 0.05$	$0.36 \pm 0.12$	$-2.14 \pm 55.86$	22633	4000
Kernel	$0.07 \pm 0.05$	$0.36 \pm 0.12$	$-10.29 \pm 222.12$	22545	3200